

Kepler's Sieve

Learning Asteroid Orbits
from Telescopic Observations

Rycroft Group Meeting: 30-Mar-2021

Michael S. Emanuel


Acknowledgments

- Masters Thesis Advisor: Pavlos Protopapas
- PhD Advisor: Chris Rycroft

Introduction

The Asteroid Search Problem

- Many asteroids (about 958,000 known) in the Solar System
- We want to learn their orbits
- Biggest data source: telescope detections
- Easy once you know which detection matches which asteroid
- This is like a jigsaw puzzle with millions of pieces!



This puzzle is too hard even with the COVID-19 Lockdown

Combining Tracklets vs. Orbital Element Search

- “Tracklet”: two detections close to each other in time and direction
- Existing search methods: greedy search over tracklets
 - Try to extrapolate a tracklet to find additional detections
 - Attempt to fit an orbit when you have enough tracklets
- Drawbacks
 - Myopic – can only connect detections made close in time
 - Suffers from combinatorial explosion
- Proposed novel method: search Orbital Elements
 - 6D space; large, but scales well
 - Cost scales as $N_{\text{ast}} \cdot N_{\text{obs}}$ rather than N_{obs}^r
 - But can we make it work?

Search Overview

- Initialize candidate orbital elements $a, e, i, \Omega, \omega, f$
- Mixture parameters: N_h, R, τ
- Compute position \mathbf{q} and velocity \mathbf{v} from candidate elements
- Compute direction \mathbf{u}_{pred} from \mathbf{q}, \mathbf{v} ; include light time and topos
- Compute distance s from \mathbf{u}_{pred} to \mathbf{u}_{obs} for ZTF observations
- Compute log likelihood \mathcal{L}_i for each candidate element
- Gradient descent...
- The rest is details! Which you will now hear all about...

Integrating the Solar System

REBOUND Integrator for N-Body Problem

- REBOUND is a modern, open source integrator
 - github.com/hannorein/rebound
- It numerically solves the gravitational N-body problem
- Considered the “gold standard” for orbits in this work
- IAS15 adaptive integrator uses Gauss-Radau quadrature and a “predictor-corrector” scheme
- Horizons: API provided by NASA JPL to obtain state vectors (position and velocity) of objects in the Solar System
- Considered “gold standard” for initial conditions of an integration

Keplerian Orbital Elements







Image Credits: Wikipedia, Cool Cosmos

- Semi-major axis a and eccentricity e describe the size and shape of the orbital ellipse
- Inclination i , ascending node Ω , perihelion ω are angles orienting orbit in the ecliptic plane
- True anomaly f is location of the body on its orbital ellipse

Validating Integration vs. Horizons

- Integrate three collections of massive bodies for 40 years at daily interval
- Initial conditions from Horizons at MJD 58600 / 2019-04-27
- Test results on first 10 IAU asteroids; query their positions from Horizons
- Report error in position (AU) and instantaneous angle from asteroid to Earth (arc seconds)
- Accuracy is excellent!
 - RMS error on planets is 5.4E-6 AU
 - Angle error from asteroids to planets 0.8 arc seconds



Bulk Integration of 733,489 Asteroids

- Download asteroid orbital elements from JPL
- Data available for 733,489
- Integrate these daily for 40 years
- Save results to disk
 - REBOUND simulation archives
 - Numpy arrays
- Job takes 4:30 on 40 CPU cores
- Writes 1.37 TB output to disk

```
# Load all the asteroid elements
ast_elt = load_ast_elt()
```



ast_elt

| | Num | Name | epoch | a | e | inc | Omega | omega | M | H | G | Ref | f |
|---------|---------|----------|---------|----------|----------|----------|----------|----------|----------|-------|------|---------|-----------|
| Num | | | | | | | | | | | | | |
| 1 | 1 | Ceres | 58600.0 | 2.769165 | 0.076009 | 0.184901 | 1.401596 | 1.284522 | 1.350398 | 3.34 | 0.12 | JPL 46 | 1.501306 |
| 2 | 2 | Pallas | 58600.0 | 2.772466 | 0.230337 | 0.608007 | 3.020817 | 5.411373 | 1.041946 | 4.13 | 0.11 | JPL 35 | 1.490912 |
| 3 | 3 | Juno | 58600.0 | 2.669150 | 0.256942 | 0.226699 | 2.964490 | 4.330836 | 0.609557 | 5.33 | 0.32 | JPL 108 | 0.996719 |
| 4 | 4 | Vesta | 58600.0 | 2.361418 | 0.088721 | 0.124647 | 1.811840 | 2.630709 | 1.673106 | 3.20 | 0.32 | JPL 34 | -4.436417 |
| 5 | 5 | Astraea | 58600.0 | 2.574249 | 0.191095 | 0.093672 | 2.470978 | 6.260280 | 4.928221 | 6.85 | 0.15 | JPL 108 | -1.738676 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1255499 | 1255499 | 2019 QG | 58600.0 | 0.822197 | 0.237862 | 0.220677 | 5.066979 | 3.770460 | 0.503214 | 21.55 | 0.15 | JPL 1 | 0.807024 |
| 1255501 | 1255501 | 2019 QL | 58600.0 | 2.722045 | 0.530676 | 0.113833 | 4.741919 | 2.351059 | 5.297173 | 19.21 | 0.15 | JPL 1 | -2.082964 |
| 1255502 | 1255502 | 2019 QQ | 58600.0 | 1.053137 | 0.389091 | 0.172121 | 5.648270 | 2.028352 | 3.266522 | 25.31 | 0.15 | JPL 1 | -3.081905 |
| 1255513 | 1255513 | 6331 P-L | 58600.0 | 2.334803 | 0.282830 | 0.141058 | 6.200287 | 0.091869 | 2.609695 | 18.50 | 0.15 | JPL 8 | 2.827595 |
| 1255514 | 1255514 | 6344 P-L | 58600.0 | 2.812944 | 0.664688 | 0.081955 | 3.199363 | 4.094863 | 2.738525 | 20.40 | 0.15 | JPL 17 | 3.032066 |

733489 rows × 19 columns

Validate Asteroid Integration vs. Horizons

- Test bulk asteroid integration on first 25 IAU asteroids
- Report position error in AU and angle error to Earth in arc seconds
- Excellent results! RMS 2.49E-6 AU and 0.45 arc seconds



Integrate Kepler Two Body Problem in TensorFlow

- Analytical solution to Kepler problem is an ellipse
- 5 of the 6 orbital elements a, e, i, Ω, ω constant
- The Mean Anomaly M is linear in time (2nd Law)

$$M(t) = M_0 + N \cdot (t - t_0)$$

- Kepler's Equation relates orbital anomalies:

$$M = E - e \sin(E)$$


Kepler's Equation

$$\tan\left(\frac{f}{2}\right) = \sqrt{\frac{1+e}{1-e}} \cdot \tan\left(\frac{E}{2}\right)$$

true to eccentric



- Convert M to E to f , then to Cartesian coordinates
- TensorFlow is fast! 5000 time points in $\sim 300 \mu \text{ sec}$
- Apply calibration $d\mathbf{q}$, $d\mathbf{v}$ to match REBOUND integration at input orbital elements

$$r(\theta) = \frac{a \cdot (1 - e^2)}{1 - e \cdot \cos(\theta - \theta_0)}$$



Predicting Directions from Positions


Right Ascension and Declination



- Fundamental plane is aligned with Earth's equator
- Intuitive, dates to ancient astronomers
- Two problems: precession (drift) and nutation (wobbles) in direction of North Pole


International Celestial Reference Frame (ICRF)

The Celestial Reference Frame Observed by Radio Waves at 24 GHz



- Modern system for RA/Dec
- Based on 232 extragalactic objects
- Addresses precession and nutation
 - Quasars don't move!
 - Not in **direction** anyway
- Amazingly accurate
 - ~2 milliarc-seconds
- Intuition: like using Polaris instead of Earth's axis for the North Pole
- Except you use 232 stars to get a highly accurate composite direction


From RA/Dec to Barycentric Mean Ecliptic



- RA/Dec is ideal for observing the stars
- But not for calculations involving orbits in the Solar System
- Inside the Solar System we want an inertial frame aligned with the ecliptic: BME
- Convert between ICRF and BME using astropy library

```
obs_icrs = astropy.SkyCoord(ra=ra, dec=dec, obstime=obstime, frame=ICRS)
obs_ecl = obs_icrs.transform_to(BarycentricMeanEcliptic)
```

Calculate Direction from Position and Velocity



- Need to remember light speed c is finite!
- Otherwise wrong by ~ 285 arc seconds

$$\mathbf{q}_{\text{rel}} = \mathbf{q}_{\text{ast}} - \mathbf{q}_{\text{earth}}$$

$$T_{\text{light}} = \|\mathbf{q}_{\text{rel}}\|/c$$

$$\Delta \mathbf{q}_{\text{ast}} = \mathbf{v}_{\text{ast}} \cdot T_{\text{light}}$$

$$\mathbf{q}_{\text{lt}} = \mathbf{q}_{\text{rel}} - \Delta \mathbf{q}_{\text{ast}}$$

$$\mathbf{u} = \mathbf{q}_{\text{lt}} / \|\mathbf{q}_{\text{lt}}\|$$

- Earth velocity doesn't matter, only asteroid velocity
- BME is an inertial frame
- Stellar aberration inapplicable, would be double counting



- Also need “topos adjustment” for observatory: Palomar Mountain, not geocenter!
- Topos adjustment worth 0-5 arc seconds on first 16 asteroids

Validating Astrometric Direction

- Check these results by comparing vs. JPL, SkyField
- Downloaded Mars at 3 hour intervals over 10 years (~29000 rows)
 - Both state vectors \mathbf{q} , \mathbf{v} and observer RA / Dec
- Computed astrometric directions from Earth to Mars
 - MSE and SkyField identical: 0.027 arc seconds
 - Both MSE and SkyField differ from JPL by 1.6 arc seconds
- Separately downloaded JPL RA/Dec on first 16 asteroids
- Compared to MSE direction calculated from integrated orbits
- RMS error: 0.873 arc seconds!

Analysis of ZTF Asteroid Detections

EDA of ZTF Detections




- ZTF: Zwicky Transient Facility; survey of northern sky at Palomar Mountain by Cal Tech
- Fast, deep survey: 3750 square degrees / hour to depth of 20.5 mag
- First light in 2017, but asteroid detections ramp up in July 2019; 7 months of data
- Enriched with machine learning pipeline that filters probable asteroid detections
- 5.69 million possible asteroid detections used for masters thesis in May 2020
- Data includes: MJD, RA, DEC, MAG
- Update: As of March 2021 now have ~158 million asteroid detections in a database!

Converting Cartesian to Angular Distance

- How far apart are two directions in the sky?
- Convert RA/Dec to directions u_1 and u_2 in the BME
- Compute Cartesian distance s between u_1 and u_2
- Angular distance θ is geodesic (great circle distance)


$$\sin(\theta/2) = s/2$$



Nearest Asteroid to Each ZTF Detection

- Compute direction $u_{\text{obs}} = (u_x, u_y, u_z)$ from RA/Dec for each detection
- Compute direction u_{ast} for every asteroid in the catalogue
- $5.7\text{E}6$ detections x $7.3\text{E}5$ asteroids = $4.2\text{E}12$ (4.2 trillion) interactions!
 - Too big for naïve brute force attack
- “Only” 97,111 different MJDs with ZTF detections
- Work in chunks of 1000 asteroids at a time, find nearest to each ZTF
- Then perform reduction operation to find globally nearest asteroid
- Still large compute job: 25 hours on 40 CPUs, 256 GB RAM server

Nearest Asteroid: 65.7% Within 2.0 Arc Seconds!



Statistical Distribution of Distance on Sphere

- What is the statistical distribution of s if we guessed directions uniformly at random?

$$s^2 = 2 \cdot (1 - z) \quad z = 1 - s^2/2$$

- This is useful parameterization because...
- “Orange Slicing Theorem” for solid angle measure:

$$d\Omega = dz \cdot d\phi$$




- Think of Z and S as random variables:

$$Z \sim \text{Unif}(-1, 1) \quad S^2 \sim \text{Unif}(0, 4)$$

- Conditional on a max (threshold) distance τ

$$S^2 | S \leq \tau \sim \text{Unif}(0, \tau^2)$$



Distribution of Nearest Asteroid Distance

- Set a threshold distance τ and define relative squared distance V

$$V = (S/\tau)^2 \quad V \sim \text{Unif}(0, 1)$$

- We have $n = 733,489$ guesses and are picking closest

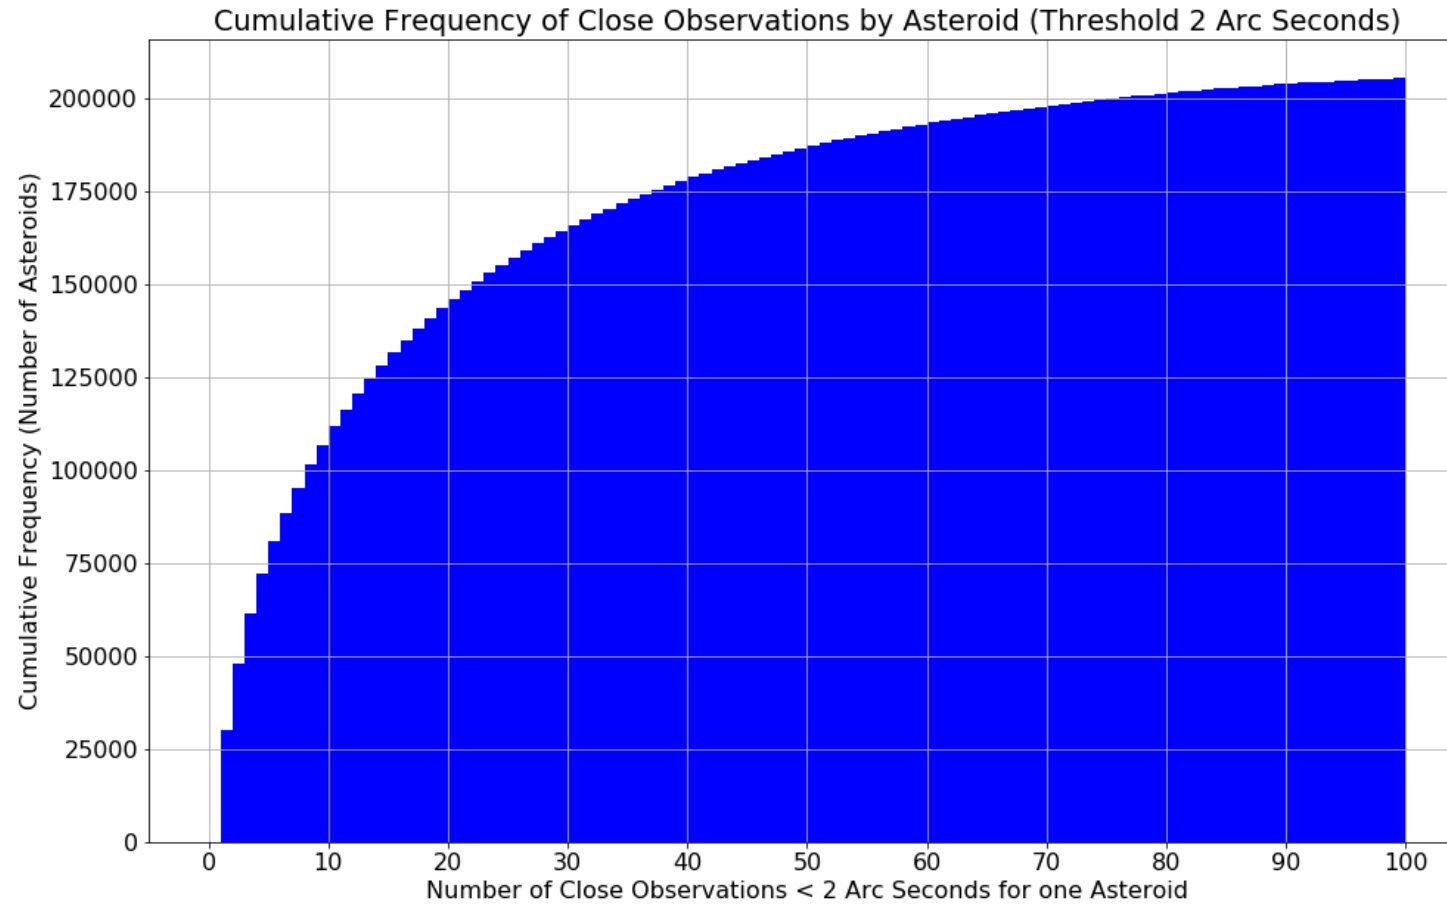
$$V_1, \dots, V_n \stackrel{i.i.d.}{\sim} \text{Unif}(0, 1)$$

- Stat 110: The minimum of n i.i.d. uniforms has a Beta distribution

$$U_{(1)} \sim \text{Beta}(1, n)$$

- How many hits at 2.0 arc seconds would we get by luck?
- Only 98. But we got 3.75 million of them!
- Conclusion: This whole apparatus works to a tolerance of 2.0 arc seconds

Cumulative Distribution of Hits per Asteroid



- Suppose we want to rebuild the asteroid catalog...
- How many asteroids do we have a sporting chance of finding?
- Count hits at 2.0 arc seconds
- How many asteroids have at least
 - 20 hits? 63,746
 - 10 hits? 100,508
- We have a shot at 13.6% of the catalogue if we require 10+ hits

Asteroid Search Using Orbital Elements

Assemble ZTF Detections Near Elements

```
# Load unperturbed element batch
ztf_elt_ast = load_ztf_batch(elts=elts_ast, thresh_deg=1.0, near_ast=False)
```

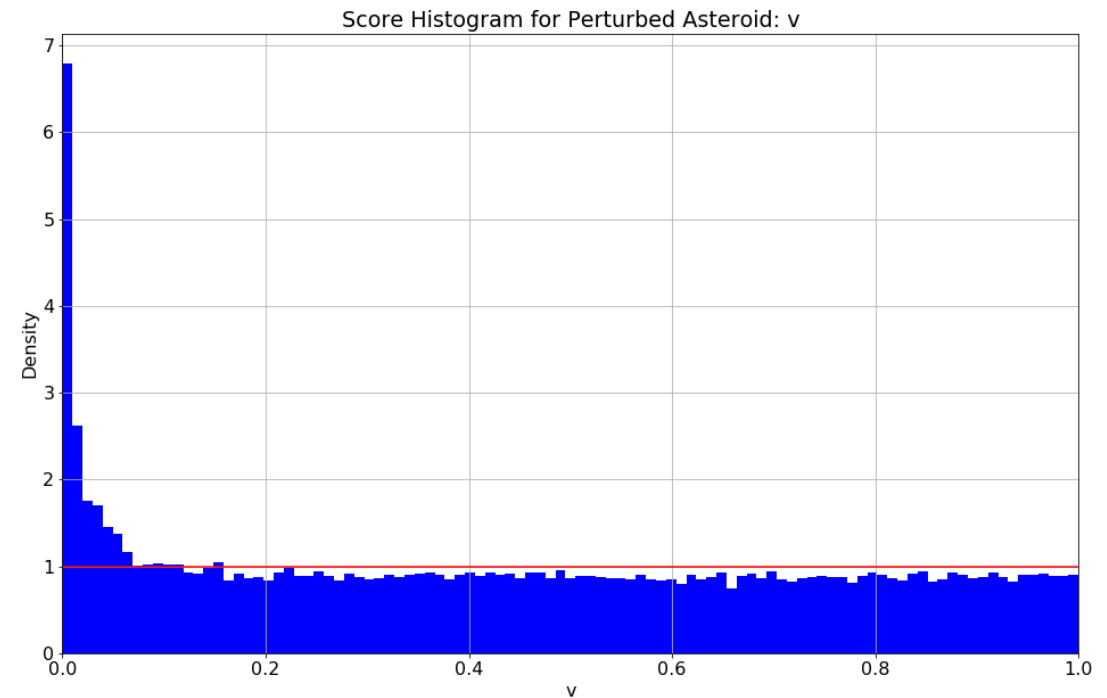

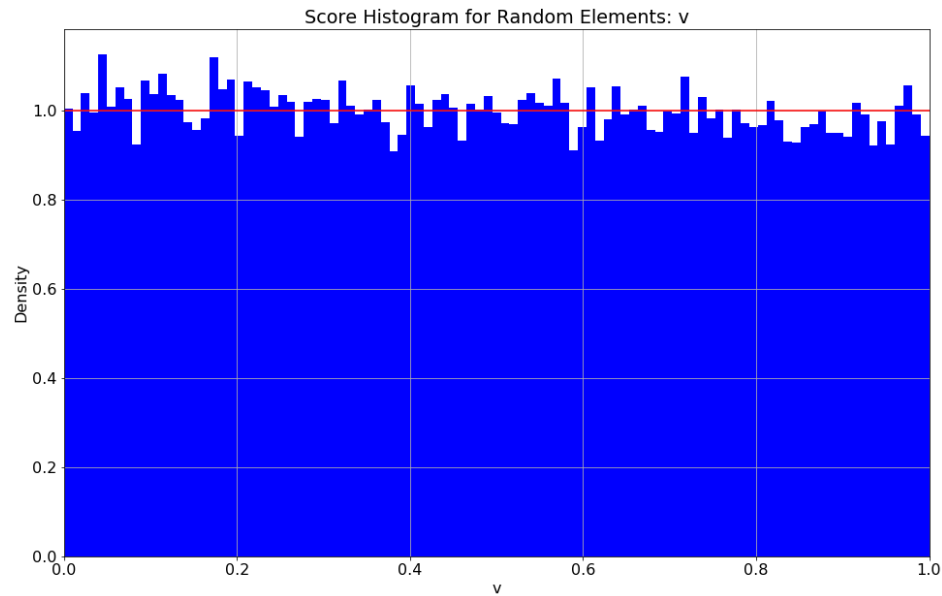
```
# Review
ztf_elt_ast[cols]
```

| | element_id | ztf_id | mjd | ra | dec | ux | uy | uz | mag_app | elt_ux | elt_uy | elt_uz | s_sec | v | is_hit |
|-------|------------|---------|--------------|------------|------------|-----------|-----------|----------|-----------|-----------|-----------|----------|-------------|----------|--------|
| 0 | 733 | 53851 | 58348.197581 | 266.229165 | -13.513802 | -0.063945 | -0.983101 | 0.171530 | 16.755600 | -0.057300 | -0.982042 | 0.179751 | 2191.408734 | 0.370552 | False |
| 1 | 733 | 73604 | 58348.197581 | 265.761024 | -13.509148 | -0.071871 | -0.982578 | 0.171389 | 16.035999 | -0.057300 | -0.982042 | 0.179751 | 3467.151428 | 0.927559 | False |
| 2 | 733 | 82343 | 58389.193252 | 270.331454 | -11.244934 | 0.005674 | -0.977422 | 0.211222 | 17.196199 | 0.000919 | -0.977996 | 0.208622 | 1124.103915 | 0.097503 | False |
| 3 | 733 | 257221 | 58685.471227 | 29.693832 | 42.180412 | 0.643725 | 0.603886 | 0.470042 | 19.289200 | 0.639004 | 0.610779 | 0.467571 | 1797.091521 | 0.249197 | False |
| 4 | 733 | 327000 | 58691.465972 | 33.104905 | 44.059131 | 0.601970 | 0.636719 | 0.481893 | 17.725201 | 0.606278 | 0.637608 | 0.475272 | 1639.539679 | 0.207419 | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 90206 | 324582 | 5650588 | 58904.176701 | 44.164238 | 29.650540 | 0.623416 | 0.752309 | 0.213037 | 18.084700 | 0.627640 | 0.750696 | 0.206212 | 1688.638104 | 0.220027 | False |
| 90207 | 324582 | 5650589 | 58904.176250 | 44.164062 | 29.650536 | 0.623417 | 0.752307 | 0.213038 | 18.165199 | 0.627641 | 0.750695 | 0.206213 | 1688.601889 | 0.220018 | False |
| 90208 | 324582 | 5650665 | 58904.176250 | 44.368640 | 28.490480 | 0.628284 | 0.753618 | 0.193182 | 19.025200 | 0.627641 | 0.750695 | 0.206213 | 2757.856412 | 0.586871 | False |
| 90209 | 324582 | 5650697 | 58904.176250 | 43.296207 | 29.505908 | 0.633424 | 0.743491 | 0.214467 | 19.852800 | 0.627641 | 0.750695 | 0.206213 | 2555.278205 | 0.503822 | False |
| 90210 | 324582 | 5650705 | 58904.176250 | 44.621045 | 29.303550 | 0.620689 | 0.756675 | 0.205398 | 19.647400 | 0.627641 | 0.750695 | 0.206213 | 1898.912116 | 0.278236 | False |

90211 rows × 15 columns

- Integrate the candidate elements on the fly in REBOUND and compute directions
- Filter the ZTF detections to those within threshold of the elements

Distribution of $V = (S/\tau)^2$ for 3 Element Batches



- Plot V for 3 batches: random, unperturbed, perturbed
- Results match theory perfectly!
- Random elements close to uniform distribution
- Unperturbed: uniform on misses with spike in first bucket
- Perturbed: in between; hits leak out to ~ 250 arc seconds

Log Likelihood Objective Function

- Mixture probability model: V mixture of h hits, $(1-h)$ misses

$$V|\text{Hit} \sim \text{Expo}(\lambda) \quad V|\text{Miss} \sim \text{Unif}(0, 1)$$

- Relate decay rate to “resolution” parameter R

$$f(v) \propto e^{-\lambda v} = e^{-\lambda s^2 / \tau^2} \quad f(v) \propto e^{-s^2 / 2R^2} \quad \lambda = \frac{\tau^2}{2R^2}$$

- The resolution R controls how tightly the model focuses

- Mixture PDF:

$$h \cdot \frac{\lambda \cdot e^{-\lambda v}}{1 - e^{-\lambda}} + (1 - h)$$

- Log Likelihood:

$$\mathcal{L}(\mathbf{v}, h, \lambda) = \sum_{j=1}^n \log \left(h_j \cdot \frac{\lambda \cdot e^{-\lambda_j v_j}}{1 - e^{-\lambda_j}} + 1 - h_j \right)$$

Search Overview

- Six trainable orbital elements $a, e, i, \Omega, \omega, f$; epoch not trainable
- Three trainable mixture parameters: N_h, R, τ
- Compute position \mathbf{q} and velocity \mathbf{v} from candidate elements
- Compute direction \mathbf{u}_{pred} from \mathbf{q}, \mathbf{v} ; include light time and topos
- Compute distance s from \mathbf{u}_{pred} to \mathbf{u}_{obs} for ZTF observations
- Compute log likelihood \mathcal{L}_i for each candidate element
- Gradient descent...
- The rest is details!

Search Techniques I: Uniform Scale, Gradient Clipping and Independent Weights

- Control variables on uniform scale in $[0, 1]$
 - e.g. $a = a_{\min} \times \exp(a \times \log(a_{\max} / a_{\min}))$; $a_{\text{trainable}}$ in $[0,1]$
- Clip gradients by norm; $\max || \text{Grad } \mathcal{L} || = 1$
 - would be better to do this elementwise, but requires custom class
- Track log likelihood and hits for each candidate element before summing them in the objective function
- Revert changes only on elements that got worse during an episode
- Weight log likelihood for each element in batch independently

$$\mathcal{L} = \sum_{i=1}^b w_i \cdot \mathcal{L}_i$$

- equivalent to controlling 64 learning rates independently
- reduce learning rate on an element when it overshoots

Search Techniques II: Mixture vs. Joint Mode, Encouraging Convergence

- Joint mode: learn all parameters jointly
- Mixture mode: only learn N_h , R , τ
- Learning rate: 2^{-12} in mixture mode vs. 2^{-16} in joint mode
- Modified objective function in mixture mode

$$\mathcal{L} = \sum_{i=1}^b w_i \cdot \frac{\mathcal{L}_i}{R_i^\alpha \cdot \tau_i^\beta}$$

- Theoretical motivation: likelihood would always look better with a larger τ ; this encourages the model to converge
- Like adjusting score for degree of difficulty in diving and gymnastics





Asteroid Search Results



Comparing Two Orbital Elements

- How far apart are two 6D orbital elements ε_1 and ε_2 ?
- A naïve Euclidean norm makes no sense at all
- Idea 1: inject the elements into space at a set of times
 - The distance between two elements is the mean distance in AU between the orbits they describe
 - Set 240 sample time points at monthly intervals from 2010 to 2030
- Idea 2: transform elements into low dimensional Cartesian space
 - Try to make each component approximately normal
 - Try to make joint distribution approximately multivariate normal
 - Use the Mahalanobis distance on these transformed elements



Train Known Asteroid Elements: Unperturbed



- Start with correct elements but uninformed mixture parameters
- Convergence is almost perfect
- Recovered Elements: 64 (100%)
- Hits: 162.6
- Resolution: 3.0 arc seconds
- Log Like: 1097





Fit Quality: Unperturbed Elements





- Do recovered elements match the nearest asteroid?
- $4.6\text{E-}8$ AU mean distance
- $5.7\text{E-}6$ covariance norm
- The fit is almost perfect
- Big deal, this is about as hard as hitting a baseball off a tee...



Train Known Asteroid Elements: Small Perturbation



- Small perturbation:
 - 1.0% to a
 - 0.25% to e
 - 0.05 degrees to i
 - 0.25 degrees to Ω, ω, f
- Convergence is very good
- Recovered Elements: 42 (65.6%)
- Hits: 117.5
- Resolution: 18.2 arc seconds
- Log Like: 798
- **Challenge: get full convergence**







Fit Quality: Small Perturbation





- Do recovered elements match the nearest asteroid?
- $2.6\text{E-}4$ AU mean distance
- 0.012 covariance norm
- This is still a very good fit on the 42 elements that have been recovered
- This is like your little league coach lobbing the ball over the plate in batting practice...

Train Known Asteroid Elements: Large Perturbation







- Small perturbation:
 - 5.0% to a
 - 1.0% to e
 - 0.25 degrees to i
 - 1.0 degrees to Ω, ω, f
- Convergence is decent
- Recovered Elements: 12 (18.8%)
- Hits: 98.2
- Resolution: 32.4 arc seconds
- Log Like: 748
- Many of these elements were perturbed so far the original is no longer even the nearest asteroid!

Fit Quality: Large Perturbation



- Do recovered elements match the nearest asteroid? Quite well.
- $4.5\text{E-}4$ AU mean distance
- 0.032 covariance norm
- This is a decent fit on the 12 elements that have been recovered
- This is a lot harder than the last task-like facing a high school pitcher...

Search Known Asteroids with Random Initializations



- ~4096 random seeds
- Each batch started with 1024 random elements
- Selected best 64 by mean $\log(v)$
- Trained for ~5 days on 4 GPUs
- Reported fits with ≥ 8 hits and resolution < 20 arc seconds
- Recovered Elements: 125
- Hits: 19.2
- Resolution: 9.3 arc seconds
- Nearest Asteroid Distance: $2.66\text{E-}3$ AU
- Nearest Asteroid Cov. Norm: 0.73
- Comments on fit quality:
 - Decent fit on some
 - Probably spurious on others
 - Overall shows that this can work
 - But not ready for production
- Baseball analogy continued: facing Roger Clemens, but trying to get hit by the pitch to get on base cheaply

Masters Thesis Conclusions

- Prove that asteroid search over orbital elements works
 - Need an adequate initialization and representation in data set
- Built a working prototype in TensorFlow
 - First demonstration of efficient astrometric computations on GPU?
- High quality integration of the Solar System and astrometric directions
 - Associated each of $5.7E6$ ZTF observations to nearest asteroid
 - $4.2E12$ interactions, possibly novel and useful data set
- Proof of concept for an automated asteroid search pipeline
 - Random initialization is not going to work, need to initialize intelligently

Progress Since June 2020

New Baby + Pandemic + House Move $\stackrel{?}{=}$ Research



- Between a new baby (Ruth, August 2020)
- and pandemic parenting
- and moving from Boston to Newton...
- Research output has been “suboptimal”
- But I have made a fair amount of progress on the unglamorous back end

From Numpy to Database Back End

- Masters Thesis used “default” strategy to manage data
 - Loose Numpy files in a data folder, organized carefully
- Efficient for prototyping, but does not scale for “big data” applications
- Started with 5.7 million detections in early snapshot
 - Painful process to upload new data
 - ZTF dataset is now up to 158 million candidate asteroid detections!
- Also, what if we want to merge data from multiple surveys, e.g. Pan-STARRS or the upcoming Vera Rubin telescope?
- Solution: put all the data into a powerful SQL database!
 - I chose a Maria-DB instance because it is fully open source

Setting up a SQL Database

- Running an instance of a SQL database is easier than you think
 - Popular variants include MySQL, MariaDB, PostgreSQL and MS SQL Server
- If you have a Linux computer, it's almost as easy as

```
$ sudo apt install mariadb-server && sudo apt install mariadb-client
```
- If you want to do serious computation, it is probably more practical to use an instance hosted by a cloud provider e.g. AWS
 - I am both crazy and a dinosaur, so I have a rack mounted server in my basement
 - This is probably not a great strategy for non crazy people
- But installing your own instance is still a great way to learn and do low cost prototyping before you start paying for big rented instances

Python Database Connections

- Extracting data from a SQL database into Python is easy
- Pandas DataFrame is familiar to most of us, very similar abstraction to a SQL resultset
 - Utility functions `sql2df`, `sp2df` convert a SQL statement or stored procedure call into a Pandas DataFrame. Very nice!
- Writing to SQL databases from Python is harder than it should be
 - Pandas comes with a canned method to this,
 - Works fine on toy data sets, but it is execrably slow on real data...
 - converting every row to be insert to a line of SQL text
 - I ended up writing a function to convert a DataFrame into CSVs in parallel with Dask; this isn't great, but it's usable

SQL Database 101 by Example

Asteroid Enter a SQL expression to filter results (use Ctrl+Space)

| | AsteroidID | AsteroidNumber | AsteroidName | BodyID | H | G |
|------|------------|----------------|--------------|-----------|------|------|
| Grid | | | | | | |
| 1 | 1 | 1 | Ceres | 1,000,001 | 3.4 | 0.12 |
| 2 | 2 | 2 | Pallas | 1,000,002 | 4.2 | 0.11 |
| 3 | 3 | 3 | Juno | 1,000,003 | 5.33 | 0.32 |
| 4 | 4 | 4 | Vesta | 1,000,004 | 3 | 0.32 |
| 5 | 5 | 5 | Astraea | 1,000,005 | 6.9 | 0.15 |
| 6 | 6 | 6 | Hebe | 1,000,006 | 5.8 | 0.24 |
| 7 | 7 | 7 | Iris | 1,000,007 | 5.6 | 0.15 |
| 8 | 8 | 8 | Flora | 1,000,008 | 6.5 | 0.28 |
| 9 | 9 | 9 | Metis | 1,000,009 | 6.3 | 0.17 |
| 10 | 10 | 10 | Hygiea | 1,000,010 | 5.5 | 0.15 |

| Column Name | # | Data Type | Not Null | Auto Increment | Key | Default | Expression | Comment |
|--------------------|---|-------------|----------|----------------|-----|---------|------------|--|
| AsteroidID | 1 | int(11) | [v] | [] | PRI | | | Internal integer ID for Asteroids; in practice this is the same as AsteroidNumber field |
| AsteroidNumber | 2 | int(11) | [v] | [] | UNI | | | The IAU asteroid number when it exists; otherwise a sequential counter starting at 1,000,000 |
| AsteroidName | 3 | varchar(32) | [v] | [] | UNI | | | The IAU asteroid name where applicable or asteroid designation |
| BodyID | 4 | int(11) | [v] | [] | UNI | | | Foreign key to KS.Body table |
| IsNumberedAsteroid | 5 | tinyint(1) | [v] | [] | | | | Flag indicating whether or not this is an IAU numbered asteroid |
| H | 6 | double | [v] | [] | | | | The H brightness parameter |
| G | 7 | double | [v] | [] | | | | The G brightness parameter |

Horizons State Vectors as DB Tables

| Grid | TimeID | HorizonsBodyID | mjd | qx | qy | qz | vx | vy | vz |
|------|------------|----------------|--------|---------|---------|--------|--------|--------|-------|
| 1 | 58,176,000 | 1 | 40,400 | 0.362 | -0.117 | -0.043 | 0.003 | 0.028 | 0.002 |
| 2 | 58,176,000 | 2 | 40,400 | 0.613 | -0.397 | -0.041 | 0.011 | 0.017 | -0 |
| 3 | 58,176,000 | 3 | 40,400 | 0.121 | -1.009 | -0 | 0.017 | 0.002 | 0 |
| 4 | 58,176,000 | 4 | 40,400 | -0.11 | -1.459 | -0.028 | 0.014 | 0 | -0 |
| 5 | 58,176,000 | 5 | 40,400 | -5.38 | -0.851 | 0.124 | 0.001 | -0.007 | 0 |
| 6 | 58,176,000 | 6 | 40,400 | 7.894 | 4.837 | -0.398 | -0.003 | 0.005 | 0 |
| 7 | 58,176,000 | 7 | 40,400 | -18.265 | -1.166 | 0.233 | 0 | -0.004 | -0 |
| 8 | 58,176,000 | 8 | 40,400 | -16.055 | -25.706 | 0.899 | 0.003 | -0.002 | -0 |
| 9 | 58,176,000 | 9 | 40,400 | -30.483 | 2.744 | 8.523 | 0 | -0.003 | 0 |
| 10 | 58,176,000 | 10 | 40,400 | 0.005 | 0.001 | -0 | -0 | 0 | -0 |

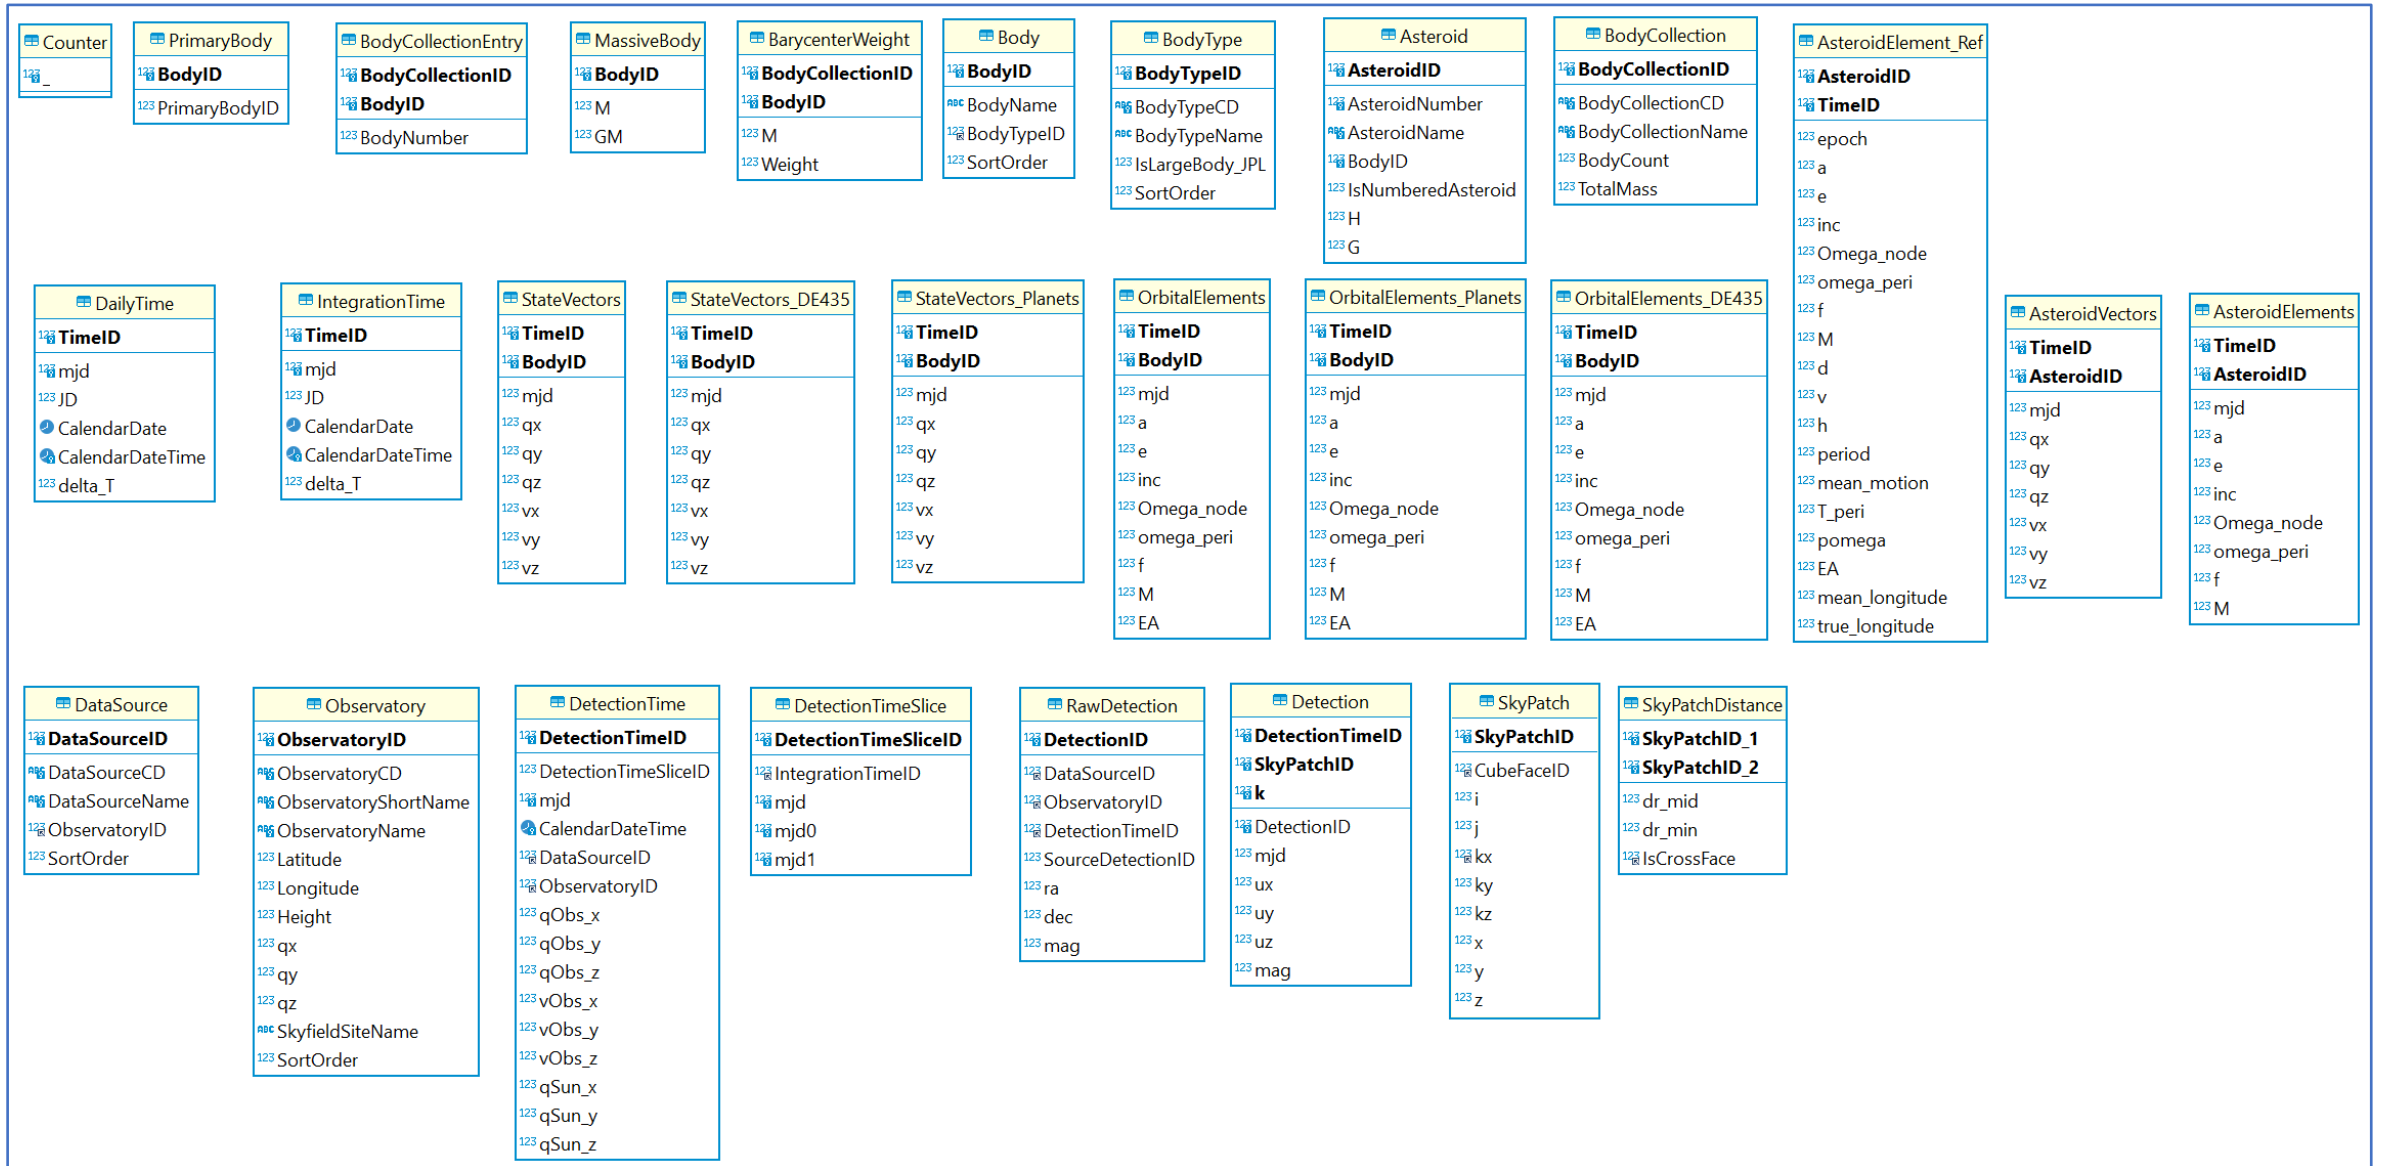
```

SELECT
  b.BodyID,
  b.BodyName,
  COALESCE(mb.M, 0.0) AS m,
  hv.qx,
  hv.qy,
  hv.qz,
  hv.vx,
  hv.vy,
  hv.vz
FROM
  JPL.HorizonsVectors AS hv
  INNER JOIN JPL.HorizonsBody AS hb ON hb.HorizonsBodyID = hv.HorizonsBodyID
  INNER JOIN JPL.HorizonsTime AS ht ON ht.TimeID = hv.TimeID
  INNER JOIN KS.Body AS b ON b.BodyID = hb.BodyID
  LEFT JOIN JPL.MassiveBody AS mb ON mb.HorizonsBodyID = hv.HorizonsBodyID
WHERE
  hb.HorizonsBodyName = HorizonsBodyName AND ht.TimeID = @TimeID;
    
```

| Grid | HorizonsBodyID | HorizonsBodyName | Body | BodyID |
|------|----------------|--------------------------|------|--------|
| 1 | 1 | LB.Mercury Barycenter | 2 | 1 |
| 2 | 2 | LB.Venus Barycenter | 2 | 2 |
| 3 | 3 | LB.Earth-Moon Barycenter | 2 | 3 |
| 4 | 4 | LB.Mars Barycenter | 2 | 4 |
| 5 | 5 | LB.Jupiter Barycenter | 2 | 5 |
| 6 | 6 | LB.Saturn Barycenter | 2 | 6 |
| 7 | 7 | LB.Uranus Barycenter | 2 | 7 |
| 8 | 8 | LB.Neptune Barycenter | 2 | 8 |
| 9 | 9 | LB.Pluto Barycenter | 2 | 9 |
| 10 | 10 | LB.Sun | 1 | 10 |
| 11 | 199 | LB.Mercury | 3 | 199 |
| 12 | 299 | LB.Venus | 3 | 299 |
| 13 | 301 | LB.Moon | 4 | 301 |
| 14 | 399 | LB.Earth | 3 | 399 |

- HorizonsBodyID is called a **primary key**
- Use integer identifiers for performance
- HorizonsVectors includes HorizonsBodyID
- This is called a **foreign key** in DB design
- Follow a simple naming rule
 - TableNameID is the PK to TableName...
- SQL query at left selects the state vectors for the sun and planets at a given date

Entity Relationship Diagram: KS



ER Diagram: JPL and ZTF

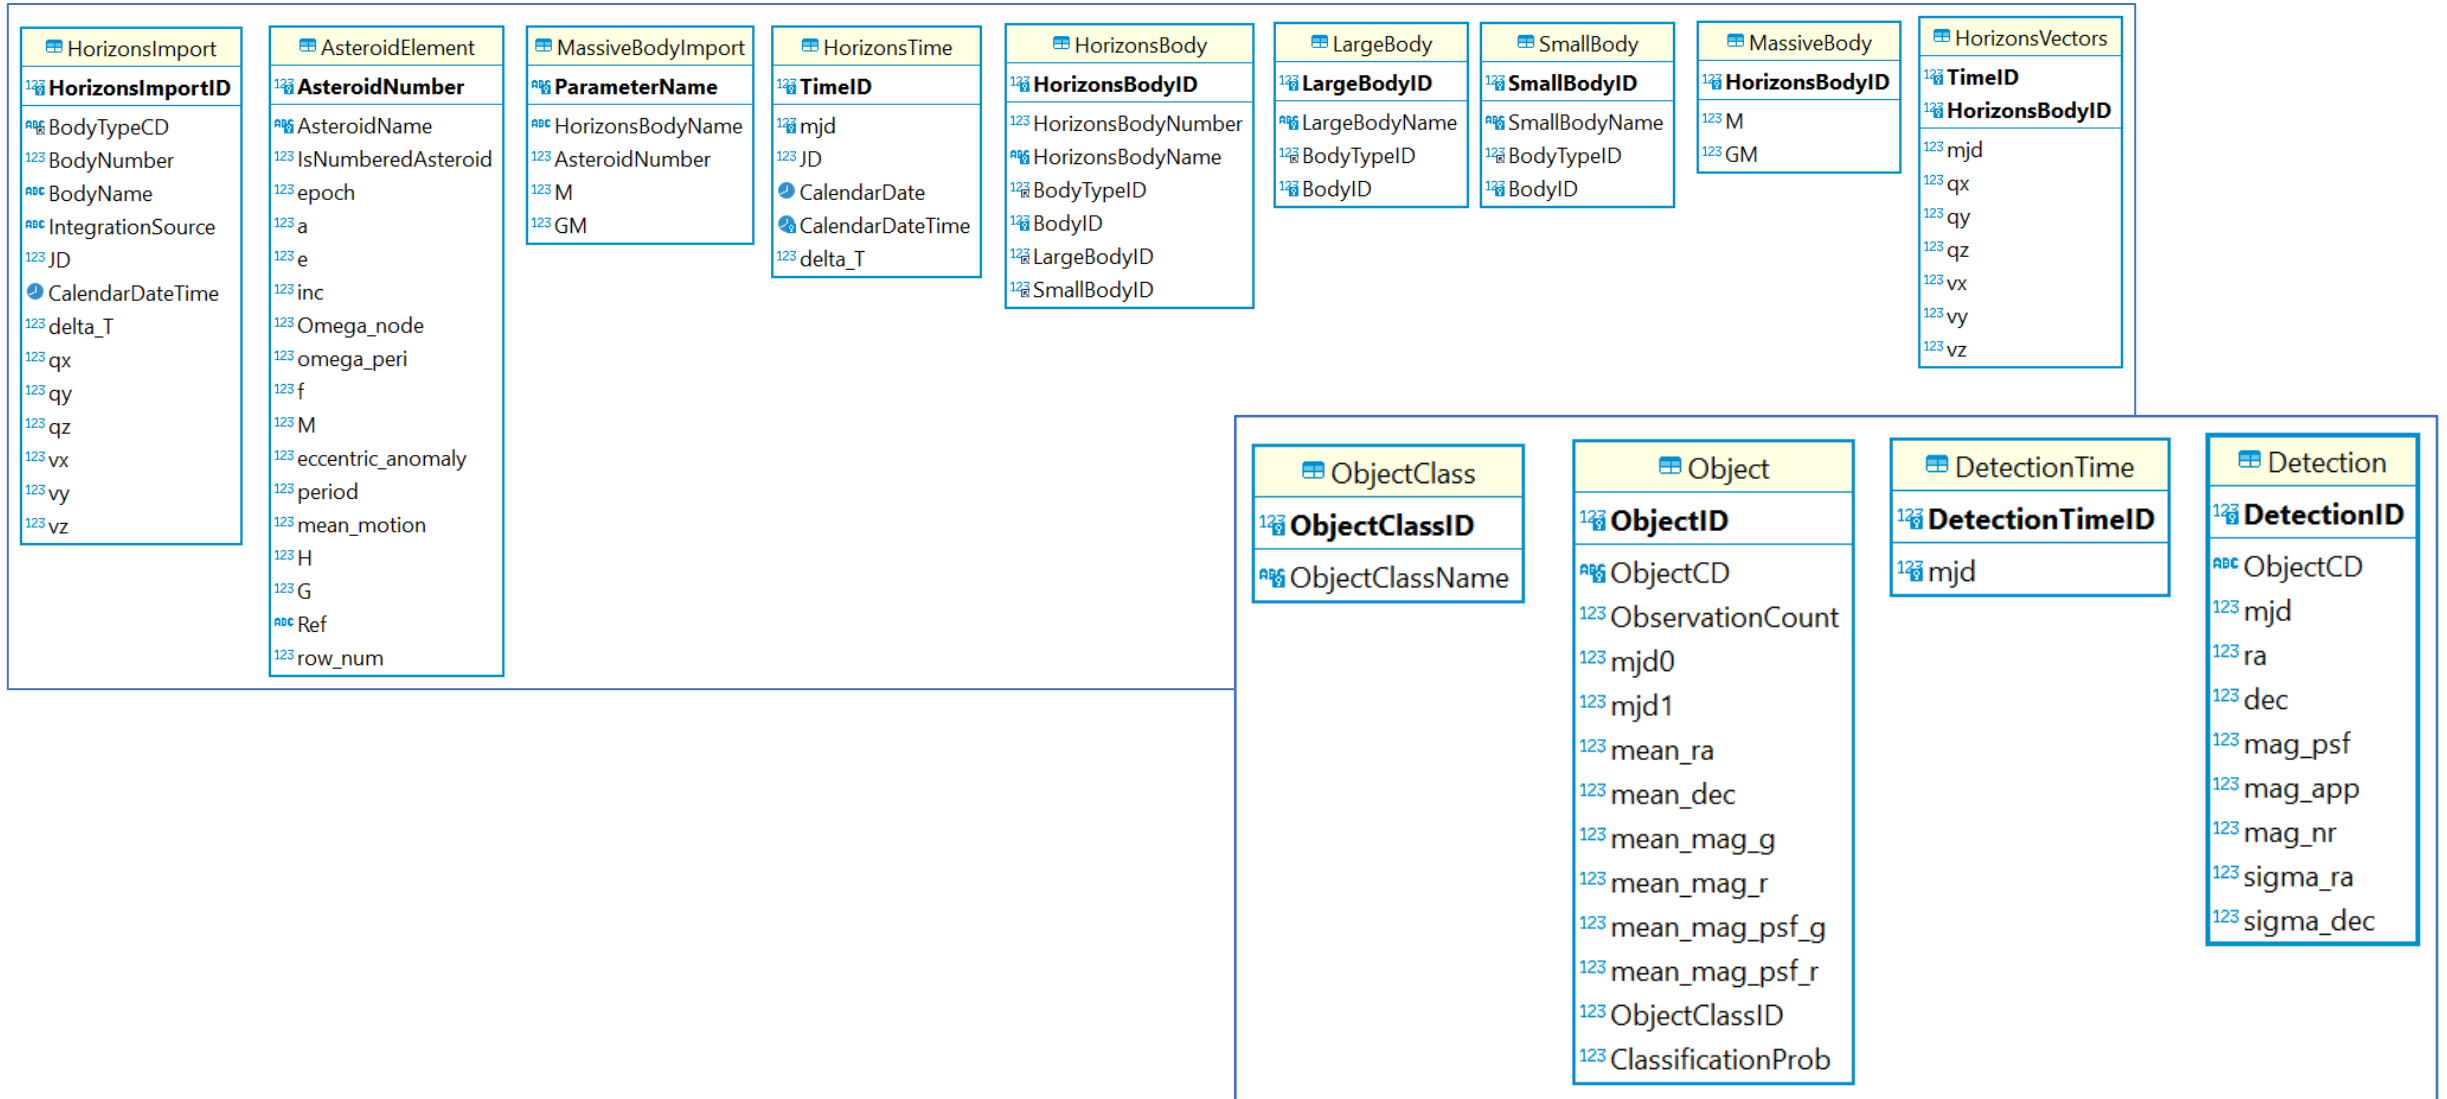




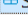

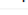








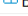

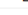









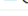











Table Descriptions and Sizes

| Table Name | Engine | Row Count | Data Length | Index length | Description |
|---|----------------------|---------------|-------------|--------------|--|
|  AsteroidElements | Aria | 3,595,918,656 | 291G | 109G | Keplerian orbital elements (a, e, i, Omega, omega, f) for all known asteroids computed in rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. |
|  AsteroidVectors | Aria | 3,595,918,656 | 263G | 109G | State vectors (position and velocity) for asteroids computed in Rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. Initial conditions for asteroids are from JPL Horizons. |
|  SkyPatchDistance | Aria | 1,861,084,200 | 57G | 74G | Distance between two SkyPatch cells; only cataloged for neighbors that are reasonably close. |
|  OrbitalElements_DE435 | Aria | 253,440,352 | 20G | 6.2G | State vectors (position and velocity) for Solar Systems bodies computed in rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. Includes record of integration errors. |
|  StateVectors_DE435 | Aria | 254,160,353 | 18G | 6.2G | State vectors (position and velocity) for Solar Systems bodies computed in rebound using all the massive bodies from the DE-435 integration with initial conditions at MJD 59000. Includes record of integration errors. |
|  SkyPatchGridDistance | Aria | 310,180,328 | 12G | 12G | Distance between two SkyPatchGrid cells; only cataloged for neighbors that are reasonably close. |
|  RawDetection | Aria | 158,843,412 | 9.7G | 5.6G | Detections of possible asteroids across multiple data sources, as quoted in RA/DEC by original source. |
|  SkyPatch | Aria | 25,165,824 | 3.6G | 1.1G | Collection of discrete patches of the sky corresponding to a cube in which the unit sphere is inscribed. |
|  OrbitalElements_Planets | Aria | 43,200,010 | 3.5G | 1.1G | State vectors (position and velocity) for Solar Systems bodies computed in rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. Includes record of integration errors. |
|  StateVectors_Planets | Aria | 47,520,011 | 3.5G | 1.2G | State vectors (position and velocity) for Solar Systems bodies computed in rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. Includes record of integration errors. |
|  SkyPatchGrid | Aria | 4,194,304 | 696M | 43M | SkyPatchGrid describes the 4N^2 square grid cells on one major face |
|  Counter | Aria | 16,777,216 | 675M | 144M | Utility table - enables a SQL query to emulate a for loop. Data type is signed integer; range is non-negative integers up to 2^24. |
|  IntegrationTime | Aria | 10,713,601 | 442M | 359M | Distinct time stamps at which MSE integrated positions of solar system bodies are available. |
|  AsteroidElement_Ref | Aria | 958,695 | 124M | 32M | Orbital elements of asteroids as of reference dates; used to initialize integrations. Primary for these elements is the Sun, NOT the Solar System Barycenter! See https://rebound.readthedocs.io/en/latest/faq.html#reference-frames |
|  CounterSigned | Aria | 2,097,151 | 84M | 29M | "Utility table - enables a SQL query to emulate a for loop. Data type is signed integer; range is non-negative integers up to 2^24."; |
|  Asteroid | Aria | 958,724 | 46M | 33M | Census of all known asteroids including reference to KS.Body table. Orbital Elements stored separately. |
|  Body | Aria | 958,809 | 38M | 17M | Solar System bodies used in the Kepler Sieve application. |
|  Detection | Aria | 676,053 | 38M | 18M | Detections of possible asteroids across multiple data sources; enriched with unit direction and SkyPatch. |
|  DetectionTime | Aria | 216,702 | 24M | 9.8M | Time at which one or more detections were made an observatory. Cartesian position and velocity includes topos adjustment. |
|  DetectionTimeSlice | Aria | 165,284 | 6.8M | 12M | Time at which one or more detections were made an observatory. Cartesian position and velocity includes topos adjustment. |
|  DailyTime | Aria | 37,201 | 1.5M | 1.3M | Distinct time stamps at sampled once per day. |
|  Minutes | Aria | 1,440 | 72K | 32K | Enumerate the 1440 minutes in a day to support linear interpolation in queries. |
|  BarycenterWeight | Aria | 438 | 32K | 24K | Weighting factors for bodies in collections. |
|  BodyCollectionEntry | Aria | 438 | 32K | 32K | Members of body collections. |
|  MassiveBody | Aria | 354 | 24K | 16K | Mass of heavy objects included in DE 435 integration, sources from technical comments. |
|  BodyCollection | Aria | 16 | 16K | 32K | Collections of bodies used in solar system integrations. |
|  BodyType | Aria | 6 | 16K | 24K | Types of Solar System bodies. |
|  CubeEdge | Aria | 12 | 16K | 24K | The 12 edges of a cube characterized by the pair of vertices sorted in ascending order. |
|  CubeEdgeVertex | Aria | 24 | 16K | 16K | Relate a cube edge to a cube vertex if the vertex is one end of the edge. |
|  CubeFace | Aria | 6 | 16K | 24K | The six faces of a cube described as the index of the constant axis and its value. |
|  CubeFaceEdge | Aria | 24 | 16K | 24K | Relate a cube face to a cube edge if the edge is on the perimeter of the face. |
|  CubeFaceNeighbor | Aria | 6 | 16K | 16K | Relate a cube face to its four neighbors based on the direction of approach. |
|  CubeFacePair | Aria | 24 | 16K | 16K | Each pair of distinct cube faces that are connected by an edge |
|  CubeVertex | Aria | 8 | 16K | 24K | The eight vertices of a cube where each coordinate is at +/- 1. |
|  DataSource | Aria | 1 | 16K | 40K | Data sources for asteroid detections. |
|  Observatory | Aria | 2 | 16K | 40K | Astronomical observatories and their position on earth. |
|  OrbitalElements | Aria | 0 | 8K | 8K | State vectors (position and velocity) for Solar Systems bodies computed in rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. |
| PrimaryBody | Aria | 0 | 8K | 8K | Default primary for each body when integrated in the Planets or DE435 collections. Almost always the Sun except that the primary of the Moon is Earth. |
| StateVectors | Aria | 0 | 8K | 8K | State vectors (position and velocity) for Solar Systems bodies computed in Rebound using the planets as massive bodies and initial conditions from DE435 at MJD 59000. Not limited to Solar Systems bodies. |

What Next?

Enrich ZTF Detection Data

- Compute direction $u_{\text{obs}} = (u_x, u_y, u_z)$ for every detection
- Assign integer SkyPatchID to every detection
- Associate every ZTF detection with nearest known asteroid
 - Populate a table keyed by (DetectionID, TimeSliceID) with payload including predicted direction and SkyPatchID
 - Join this against ZTF detections with match on TimeSliceID and SkyPatchID matching only on nearby parts of the sky
- Compute the exact position for every candidate asteroid that might be the nearest to each detection, then take the closest one

Generate Candidate Elements from Tracklets

- A set of orbital elements has 6 degrees of freedom
- One detection has 2 D.O.F., RA and DEC; a tracklet has 4
 - In theory a set of three detections close together has 6, but in practice you only get 4, RA, DEC, and their time derivatives
 - If you have a third detection or tracklet that is separated in time, you should be able to fit orbital elements if it they are consistent with a realistic orbit
- To discover new asteroids, probably need some way to generate a 2D space of candidate orbital elements consistent with a single tracklet
- Current idea: provisionally assign r (distance from sun to object) and r' ; these follow a well behaved distribution and are easy to work with

Mille Grazie: Thank you for Your Attention!

- Questions?
- Comments?
- Suggestions?