

Dissertation Advisors:

Professor Pavlos Protopapas
Professor Christopher H. Rycroft

Author:

Michael S. Emanuel

**Kepler’s Sieve:
Learning Asteroid Orbits from Telescopic Observations**

Abstract

A novel method is presented to learn the orbits of asteroids from a large data set of telescopic observations. The problem is formulated as a search over the six dimensional space of Keplerian orbital elements. Candidate orbital elements are initialized randomly. An objective function is formulated based on log likelihood that rewards candidate elements for getting very close to a fraction of the observed directions. The candidate elements and the parameters describing the mixture distribution are jointly optimized using gradient descent. Computations are performed quickly and efficiently on GPUs using the TensorFlow library.

The methodology of predicting the directions of telescopic detections is validated by demonstrating that out of approximately 5.69 million observations from the ZTF dataset, 3.75 million (65.71%) fall within 2.0 arc seconds of the predicted directions of known asteroids. The search process is validated on known asteroids by demonstrating the successful recovery of their orbital elements after initialization at perturbed values. A search is run on observations that do not match any known asteroids. I present orbital elements for 12 new, previously unknown asteroids with at least 8 hits within 10 arc seconds on ZTF asteroid detections.

All code for this project is publicly available on GitHub at github.com/memanuel/kepler-sieve.

Chapter 1

Asteroid Search Results

1.1 Comparing Candidate Elements to the Nearest Asteroid

Before we review the results of our asteroid search experiments, it will be helpful to have in hand a notion of how closely two sets of orbital elements match. In particular, we will test below whether or not we successfully recovered orbital elements when we started with them as an initial guess. Answering this question requires that we have a useful metric on the space of orbital elements.

I spent a fair amount of time trying to develop such a metric. While orbital elements are convenient for intuition and calculating orbits, there isn't an obvious distance metric we can put on them that makes a lot of sense. Eventually I decided that the canonical way to compare two orbits by comparing the predicted vector of positions on a set of representative dates. Logically this is hard to argue with, but it is somewhat computationally expensive compared to a computation that can be run directly on the elements.

The method `nearest_ast` searches the asteroid catalogue for the known asteroid whose orbit is closest to that predicted by the candidate elements. It delegates its work to the function `nearest_ast_elt_cart`, which is defined in `nearest_asteroid.py`. This function creates a set of 240 sample time points over 20 years spanning 2010 to 2030 sampled monthly. The resulting table of positions for the asteroid catalog is fairly large, with a size of $[733490, 240, 3]$ (5.28E8 elements and about 2.11 GB using 32 bit floats). Computing the nearest asteroid against 64 candidate elements by brute force in TensorFlow would necessitate creating a tensor with 3.38E10 elements or 135 GB of memory—two orders of magnitude too large for a high quality consumer

grade GPU with ~ 10 GB of memory. The nearest asteroid method is therefore forced to iterate through the elements one at a time, taking the norm of the difference against the table. In one important optimization, the tensor of known asteroid positions is loaded into memory once as a TensorFlow constant to avoid recomputing it every time the function is called.

I also sought to develop a sensible metric of the distance between a pair of arbitrary orbital elements. This is implemented in the same module with the function `elt_q_norm` and `nearest_ast_elt_cov`. The idea is to transform the elements to a Cartesian representation where they have a well behaved covariance matrix. In particular, the goal is to find a deterministic transform of the elements that is distributed approximately as a multivariate normal. Then the [Mahalanobis distance](#) is a natural metric on the transformed elements. This process can be reviewed in the Jupyter notebook `11_nearest_asteroid.ipynb`. I initially tried working with the full empirical distribution to convert every orbital element to a percentile and then to a normally distributed z score, but the results didn't make any sense, so I dropped that approach.

Instead, I switched to a simpler approach and standardized variables so they would have mean zero and variance 1. I attempt to make them close to normal if possible, but without overfitting against the empirical distribution. I standardized the log of the semimajor axis a and directly standardized the eccentricity e . (Even though this admits mappings from z to eccentricities outside $[0, 1]$, the mapping is only used in the direction from a reported eccentricity e to a transformed e_z that is approximately normal). The quantity $\sin(i)$ was also standardized. Here is a summary of the mathematical transformations to create approximately normal variables from a , e and i :

$$\begin{aligned} a_z &= \frac{\log(a) - E[\log(a)]}{\sqrt{\text{Var}[\log(a)]}} \\ e_z &= \frac{e - E[e]}{\sqrt{\text{Var}[e]}} \\ i_z &= \frac{\sin(i) - E[\sin(i)]}{\sqrt{\text{Var}[\sin(i)]}} \end{aligned}$$

The expectation and variance here are estimated using the sample mean and sample variance respectively.

Here are visualizations of comparing the hypothetical and empirical distributions of a and e :

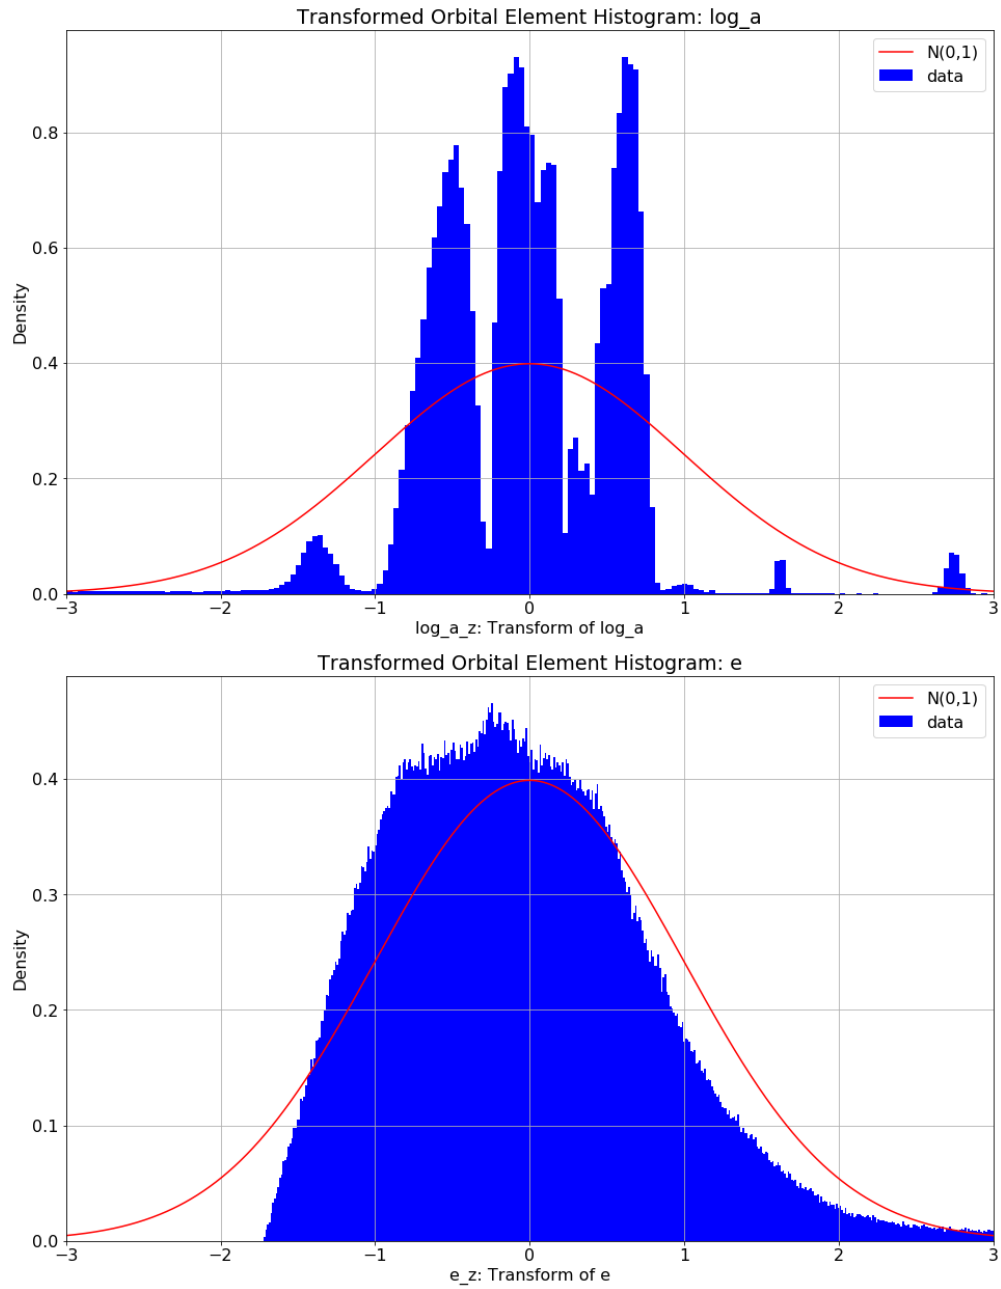


Figure 1.1: Transformations of a and e to standardized (and “approximately” normal) variables a_z and e_z .

The other three angular orbital elements Ω , ω and f are handled identically. We can inject i into Cartesian space with only its sine because it is constrained to $[-\pi, \pi]$. But the other three angles are unconstrained. I will take Ω as an example. I transform Ω into *two* variables, named `cos_Omega_z` and `sin_Omega_z`. These are not transformed empirically, but using a theoretical distribution. Let x be the sine or cosine of one of Ω , ω or f . x is mapped to a variable z that is distributed approximately normal by applying the transformation

$$u = \frac{1/2 + \arcsin(x)}{\pi}$$

$$z = \Phi^{-1}(u)$$

where Φ is normal CDF.

Below are visualizations of comparing the hypothetical and empirical distributions of $\sin(i)$ and $\sin(\Omega)$:

Better results for e and i could be obtained by using the Beta distributions noted above, for this purpose the simple standardization is adequate. The plot shown for the tranform of $\sin(\Omega)$ shows that it is very close to the theoretical distribution. I generated analogous plots for the sin and cos of Ω , ω and f which are in the Jupyter notebook. They are qualitatively similar to this one and all show excellent fits.

I have now given a recipe with which six orbital elements can be injected into R^9 . Let X be the $N \times 9$ matrix of transformed elements ($N = 733,489$ is the number of asteroids). The orbital elements are only very lightly correlated with each other, and so are the X_j except for the tightly correlated pairs with the sin and cos of the same angle. Next, using the Spectral Theorem, I find a 9×9 matrix β such that the covariance matrix of $X\beta$ (which also has shape $N \times 9$) is the 9×9 identity matrix. The only wrinkle is that I assign importance weights to the 9 columns before building X and computing β . The importance weights are:

- 1.0 for a_z and e_z
- 0.5 for i_z
- 0.1 for the sin and cos of Ω , ω and f

These are admittedly qualitative judgments on my part. I initially only compensated for the

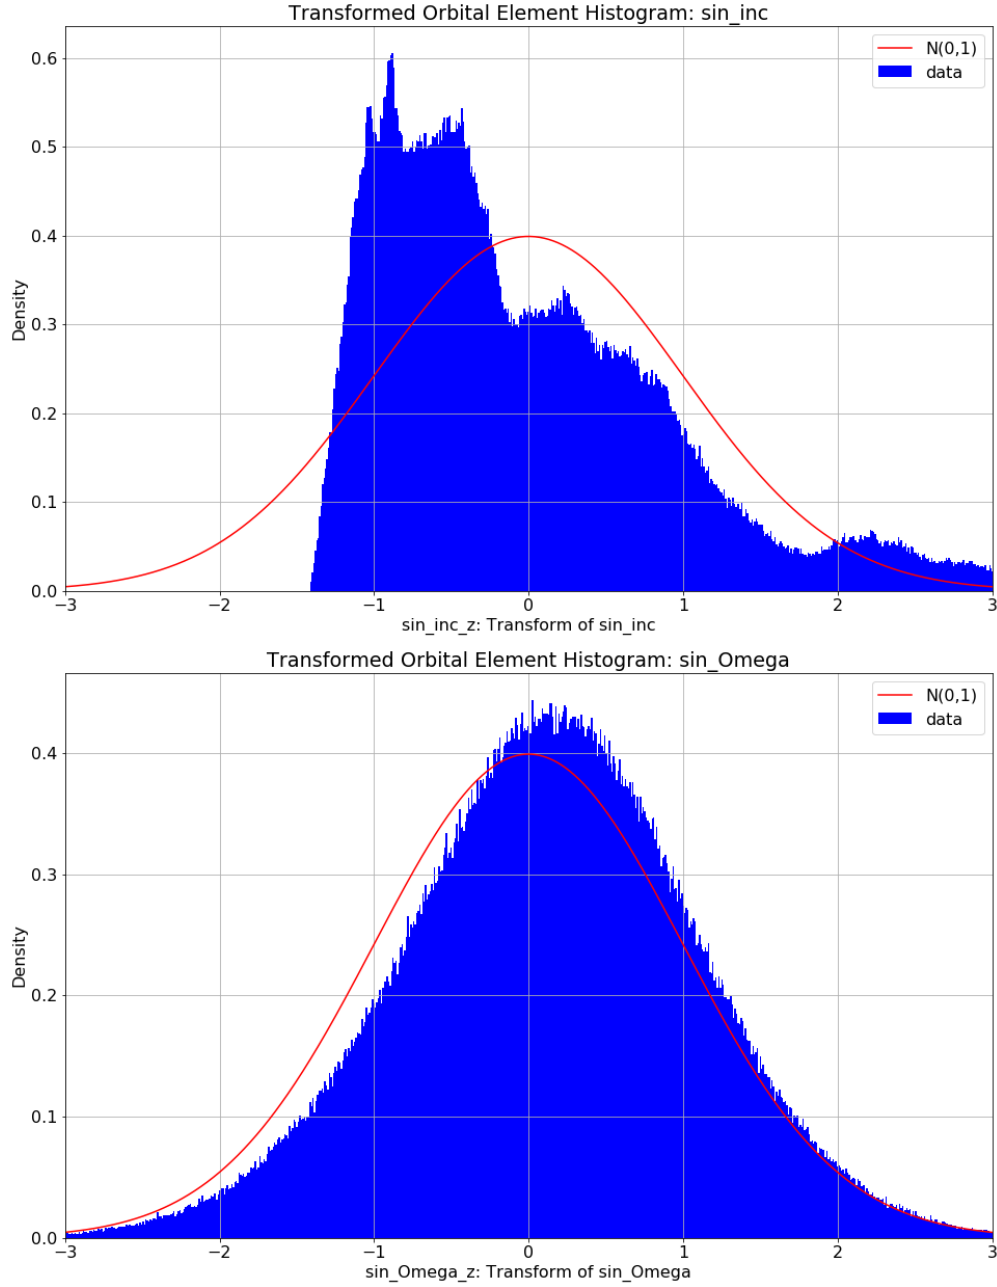


Figure 1.2: Transformations of a and e to standardized (and “approximately” normal) variables a_z and e_z .

double counting of Ω , ω and f , but I noticed that relatively small differences in e.g. ω on a near circular orbit that hardly effected the shape of an orbit were having a disproportionately large influence of covariance score.

The covariance metric between two sets of orbital elements ϵ_1 and ϵ_2 is defined by

$$\|\epsilon_2 - \epsilon_1\|_{\text{cov}} = \|\epsilon_2\beta - \epsilon_1\beta\|$$

The importance weights are rescaled so the diagonal of the covariance matrix sums to 1 and a random pair of elements should have distance 1. These calculations are also in `nearest_asteroid.py` and done by the functions `elts_to_X_cov`, `calc_beta`, `elt_q_norm` and `nearest_ast_elt_cov`. Now that we know what it means for two orbital elements to be “close,” we are ready for our first test: recovering unperturbed elements.

1.2 Recovering the Unperturbed Elements of Known Asteroids

The first and easiest proof of concept for the search process is to see if the mixture parameters will converge correctly when the search is initialized with correct orbital elements for asteroids that are well represented in the data, but with “neutral” or uninformative mixture parameters. I liken this test to a kid learning to swing a bat by trying to ball sitting on a tee. This test is demonstrated in the Jupyter notebook `14_asteroid_search_unperturbed.ipynb`. It was developed before the automated sieving routine, so it includes lower level calls to the `adaptive_search` method.

It’s worthwhile to follow through the steps to assemble the data to get familiar with how everything fits together. These two lines of codes load the ZTF data observations associated with the nearest asteroid, and count hits by asteroid number:

```
ast_elt = load_ztf_nearest_ast()
ast_num, hit_count = calc_hit_freq(ztf=ast_elt, thresh_Sec=2.0)
```

The next few lines sort the asteroids in descending order by number of hits, and assemble a data frame of the orbital elements belonging to the 64 “best” asteroids. The function `asteroid_elts` in `candidate_elements` provides a batch of candidate orbital elements that exactly match known asteroids; it assigns an `element_id` matching the original asteroid number to make it easy to check later if the fitted elements match the original. The function `load_ztf_batch` assembles the batch of ZTF observations within a threshold, here 2.0 degrees, of these candidate elements

```
elts_ast = asteroid_elts(ast_nums=ast_num_best[0:64])
ztf_elt = load_ztf_batch(elts=elts_ast, thresh_deg=2.0)
```

Reviewing the `ztf_elt` on screen we can see that there are 322,914 rows which include 10,333 hits: an average of 161.5 hits per candidate element and 3.2% of the total rows of data. A call to `score_by_elt` computes the t-score described earlier based on the mean and standard deviation of $\log(v)$. This shows a mean t-score of +45.0, which is off the charts good. It's interesting to see that a set of observations with 3.2% hits and 96.8% noise achieves such a good score. This also puts into context the challenge of the search problem: we have an average of 5045 detections within 2.0 degrees of each set of candidate elements, of which 160 are hits and the remaining 4885 are random detections belonging to other asteroids. If we want a search process to detect asteroids with as few as 8 hits in the data, we will need a process selective enough to pick out just 0.16% of the observations.

To initiate the search, we also need to choose our initial mixture parameters. we set `num_hits` to 10 and the resolution to 0.5 degrees

```
elts_add_mixture_params(  
    elts=elts, num_hits=num_hits,  
    R_deg=R_deg, thresh_deg=thresh_deg)
```

Now that we have the candidate elements and the ZTF data frame, we are ready to instantiate the asteroid search model:

```
model = AsteroidSearchModel(  
    elts=elts_ast, ztf_elt=ztf_elt,  
    site_name='palomar', thresh_deg=2.0)
```

Before we start training the model, we can get a plain text report or a visualization of the starting point. I will omit these here. The report shows that at the start of training, 10 elements are “good” with 5 or more hits, and the overall mean log likelihood is 3.13 and the mean number of hits is 3.19. ¹

¹ As I write this, I realize that something is going wrong somewhere. The initial model should show the same number of hits, 10,333, as the ZTF data frame. There must be some slippage between the CPU / `numpy` calculations and the TensorFlow GPU model on the order of 10 arc seconds or more to explain this. I will investigate this later.

Here is an excerpt from the plain text model report:

```
model.report()
```

```
Good elements (hits >= 5): 64.00
```

```
      \ log_like : hits :      R_sec : thresh_sec
```

```
Mean Good: 1057.30 : 136.75 :      7.05 : 347.11
```

```
Mean Bad :      nan :      nan :      nan :      nan
```

```
Min      : 848.14 : 15.00 :      1.97 : 112.80
```

```
Max      : 1307.20 : 190.00 :      14.87 : 697.68
```

```
Trained for 9536 batches over 149 epochs and 59 episodes (elapsed time 425 seconds).
```

I've developed a number of visualizations to assess training progress.

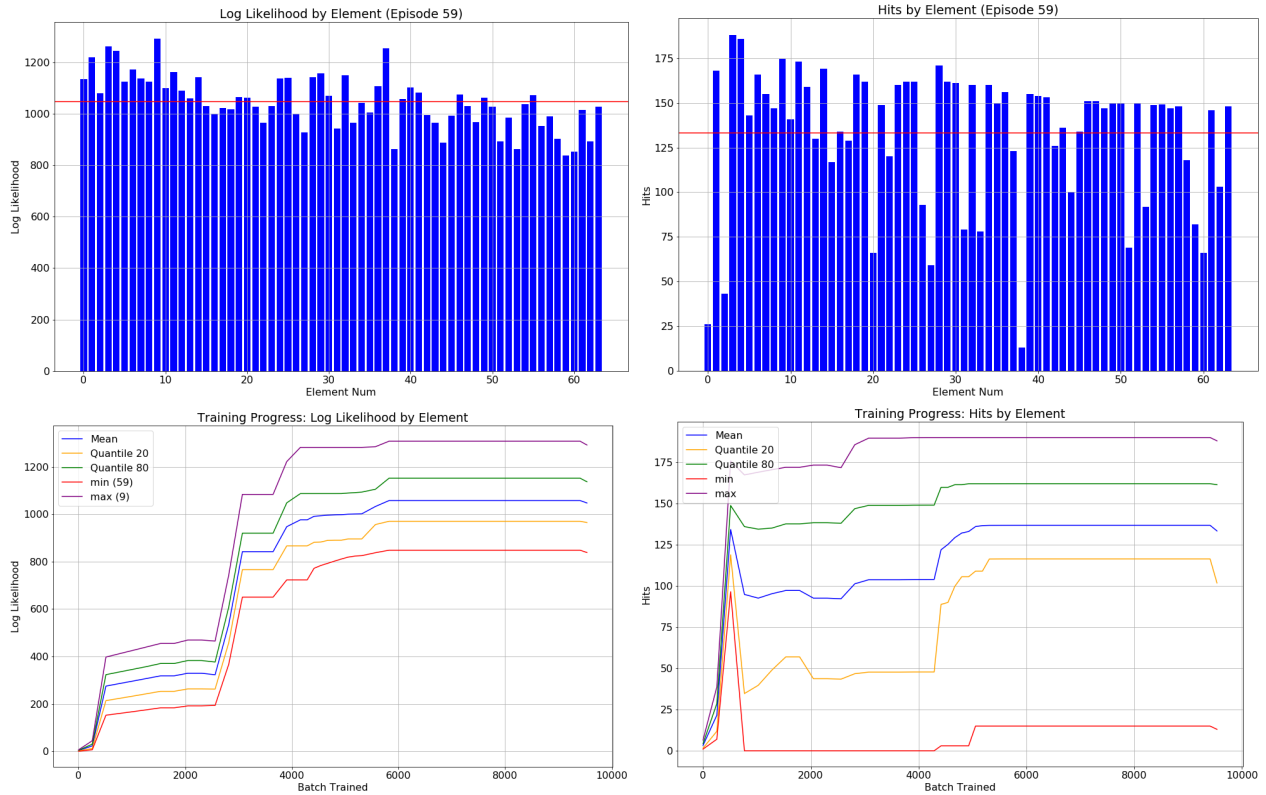


Figure 1.3: Training progress on 64 unperturbed orbital elements.

We can see that the model is behaving as hoped. It is gradually ratcheting the resolution parameter and scoring a high log likelihood as it does so. It does this without getting deked and polluting the originally correct orbital elements.

The diagnostics presented above work equally well for any set of candidate orbital elements, whether or not they are ostensibly associated with a known asteroid. In this case, we can further validate the results by comparing our fitted elements to the nearest asteroid using the two metrics




Figure 1.4: The resolution R decreases monotonically when training the unperturbed elements.

described in the previous section.

Here is a data frame comparing the recovered elements with the nearest asteroid. The simplest test is how many of 64 recovered elements have as their nearest asteroid the same asteroid used to initialize the elements. The answer is 64: the fitting process “tried” to converge back to the right asteroid every time. A more substantive question is how close did it come. Here are the geometric mean differences of two metrics:

- Distance in AU: $6.61\text{E-}6$
- Covariance Norm of elements: $2.00\text{E-}3$

This is an excellent level of agreement. The covariance norm is a good summary statistic, but may be hard to relate to astronomy. The mean absolute error in the recovered a is $4.3\text{E-}4$ and in the recovered e is $1.6\text{E-}4$.

Here are two visualizations showing the distance in AU and covariance norm to the nearest (original) asteroid for all 64 candidate elements.

	element_id	nearest_ast_num	nearest_ast_name	nearest_ast_dist	nearest_ast_q_norm
0	51921	51921	2001 QU90	0.000089	0.001428
1	59244	59244	1999 CG6	0.000049	0.003178
2	15786	15786	1993 RS	0.000022	0.000617
3	3904	3904	Honda	0.000023	0.000552
4	142999	142999	2002 VZ98	0.000091	0.001302
...
59	11952	11952	1994 AM3	0.000437	0.007174
60	134815	134815	2000 FA30	0.000117	0.001158
61	27860	27860	1995 BV2	0.000137	0.006892
62	85937	85937	1999 DL1	0.000240	0.009916
63	72911	72911	2001 OC32	0.000123	0.001776

Figure 1.5: *The nearest asteroid to the recovered elements initialized with unperturbed asteroid elements.*




Figure 1.6: Two metrics comparing the recovered orbital elements to the true elements of the asteroid in question. Both charts are plotted on a log scale with precision (reciprocal of the error) on the y axis. The geometric mean error is shown in red: $6.61\text{E-}6$ AU and $2.00\text{E-}3$ on the covariance norm.

1.3 Recovering the Perturbed Elements of Known Asteroids

1.3.1 Small Perturbation

The next experiment is similar to the previous one. This time we will apply a small perturbation to the orbital elements in our initial guess. If the last experiment was like hitting a tee ball, this one may be likened to hitting a ball gently pitched by your little league coach in batting practice. The elements are perturbed using the function `perturb_elts` in `candidate_elements.py`. The perturbation adds normally distributed random noise with the specified standard deviation to $\log(a)$, $\log(e)$, and the four angles i , Ω , ω and f . The small perturbation shifts $\log(a)$ by 0.01, $\log(e)$ by 0.0025, i by 0.05 degrees, and the the other angles by 0.25 degrees. A random seed is used for reproducible results The code to do this is

```
elts_pert= perturb_elts(elts_ast, sigma_a=0.01, sigma_e=0.0025,
                        sigma_inc_deg=0.05, sigma_f_deg=0.25,
                        sigma_Omega_deg=0.25, sigma_omega_deg=0.25,
                        random_seed=42)
```

Last time we the summary statistic based on $\log(v)$ showed a very positive t score. This time the t score has dropped to +3.71, and the model has zero hits before it begins training. Even this small perturbation is enough that the model is going to have to work quite a bit to recover the elements.

Here is the text report after sieving:

```
Good elements (hits >= 5): 32.00
      \ log_like : hits :      R_sec : thresh_sec
Mean Good:  798.22 : 73.00 :    42.71 : 1013.88
Mean Bad :   151.47 :  0.16 :   246.08 : 2130.53
Min      :    1.33 :  0.00 :    3.53 :  233.61
Max      :  1226.84 : 163.00 :  1168.61 : 2400.00
Trained for 12096 batches over 189 epochs and 71 episodes (elapsed time 520 seconds).
```

Here are summary statistics for the run on the small perturbation of real asteroid elements:

- Successfully converged for 32 out of 64 candidate elements
- Mean hits on converged elements: 73.00
- Resolution on converged elements: 42.7 arc seconds
- Distance in AU to nearest asteroid: 4.00E-4
- Covariance Norm to nearest asteroid: 1.69E-2

These results are not as strong as on the unperturbed elements, but the method is still clearly working. It's covered on half of the candidate orbital elements. The converged elements are fit well, average 73 hits at 43 arc seconds. The distance to the nearest asteroid is $4.0\text{E-}4$ AU, which is still very close and an excellent description of the orbit.

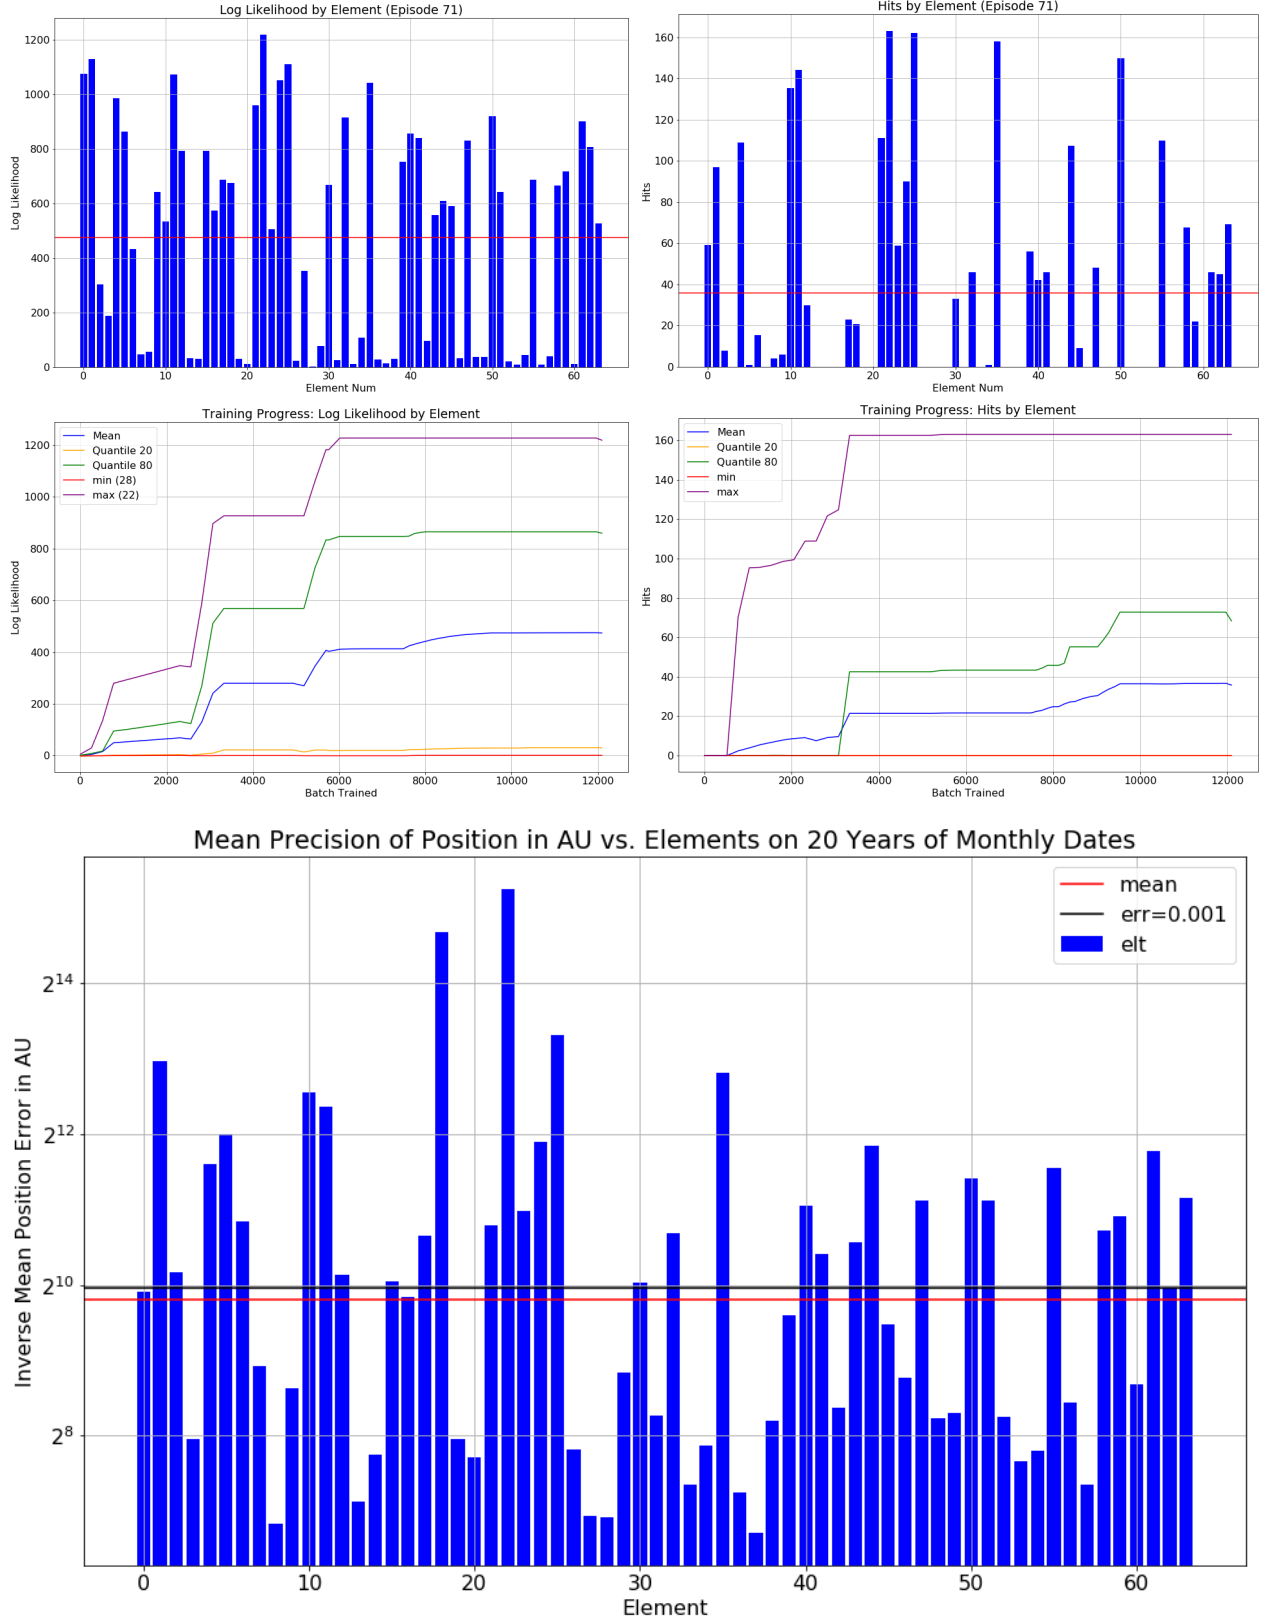


Figure 1.7: Training progress on 64 orbital elements initialized with small perturbations from real asteroids. 32 of the 64 candidate elements converge, averaging 73 hits each.

1.3.2 Large Perturbation

In our third test, we will again start with perturbed orbital elements. But this time, we will apply a much larger perturbation. This is much a much harder task. To continue with the baseball analogy, it might be like batting in a high school game. The perturbation size this time is 0.05 on $\log(a)$, 0.01 on $\log(e)$, 0.25 degrees on i , and 1.0 degree on the other three angles. While this might not sound like much at first, they are large perturbations. In fact, they are so large they led me down a painful rabbit hole. I repeatedly failed to recover the orbital elements of the original asteroids before I realized that the perturbations were large enough that in many cases, the nearest asteroid to the perturbed element was no longer the original asteroid! The results started to make much more sense when I compared each fitted element to the nearest real asteroid, regardless of whether this matched the original source of the elements before perturbation.

Here is the text report after sieving:

```
Good elements (hits >= 5):    9.00
      \ log_like : hits :    R_sec : thresh_sec
Mean Good:   950.21 :  87.33 :    26.00 :   801.43
Mean Bad :    42.19 :   0.11 :   252.31 :  2347.92
Min      :     8.83 :   0.00 :     2.92 :   341.62
Max      :  1351.26 :  163.00 :   742.89 :  2400.00
Trained for 13056 batches over 204 epochs and 76 episodes (elapsed time 547 seconds).
```

Here are summary statistics for the run on the small perturbation of real asteroid elements:

- Successfully converged for 9 out of 64 candidate elements
- Mean hits on converged elements: 87.00
- Resolution on converged elements: 26.0 arc seconds
- Distance in AU to nearest asteroid: 4.11E-4
- Covariance Norm to nearest asteroid: 3.18E-2

This time we've only converged on 9 of the 64 orbital elements. But the encouraging news is that when we have converged, the fit is as good as it was before. The average hits are 87 and the resolution is 26.0 arc seconds. The distance to the nearest asteroid is comparable to the batch initialized with small perturbations, at 4.0E-4. This is telling us something important: when the model starts from a good enough guess that it has a path in search space to the local maximum, it

will converge to a good solution. Starting from a poor initialization will reduce the probability of successful convergence, but it doesn't dilute the quality of the results.

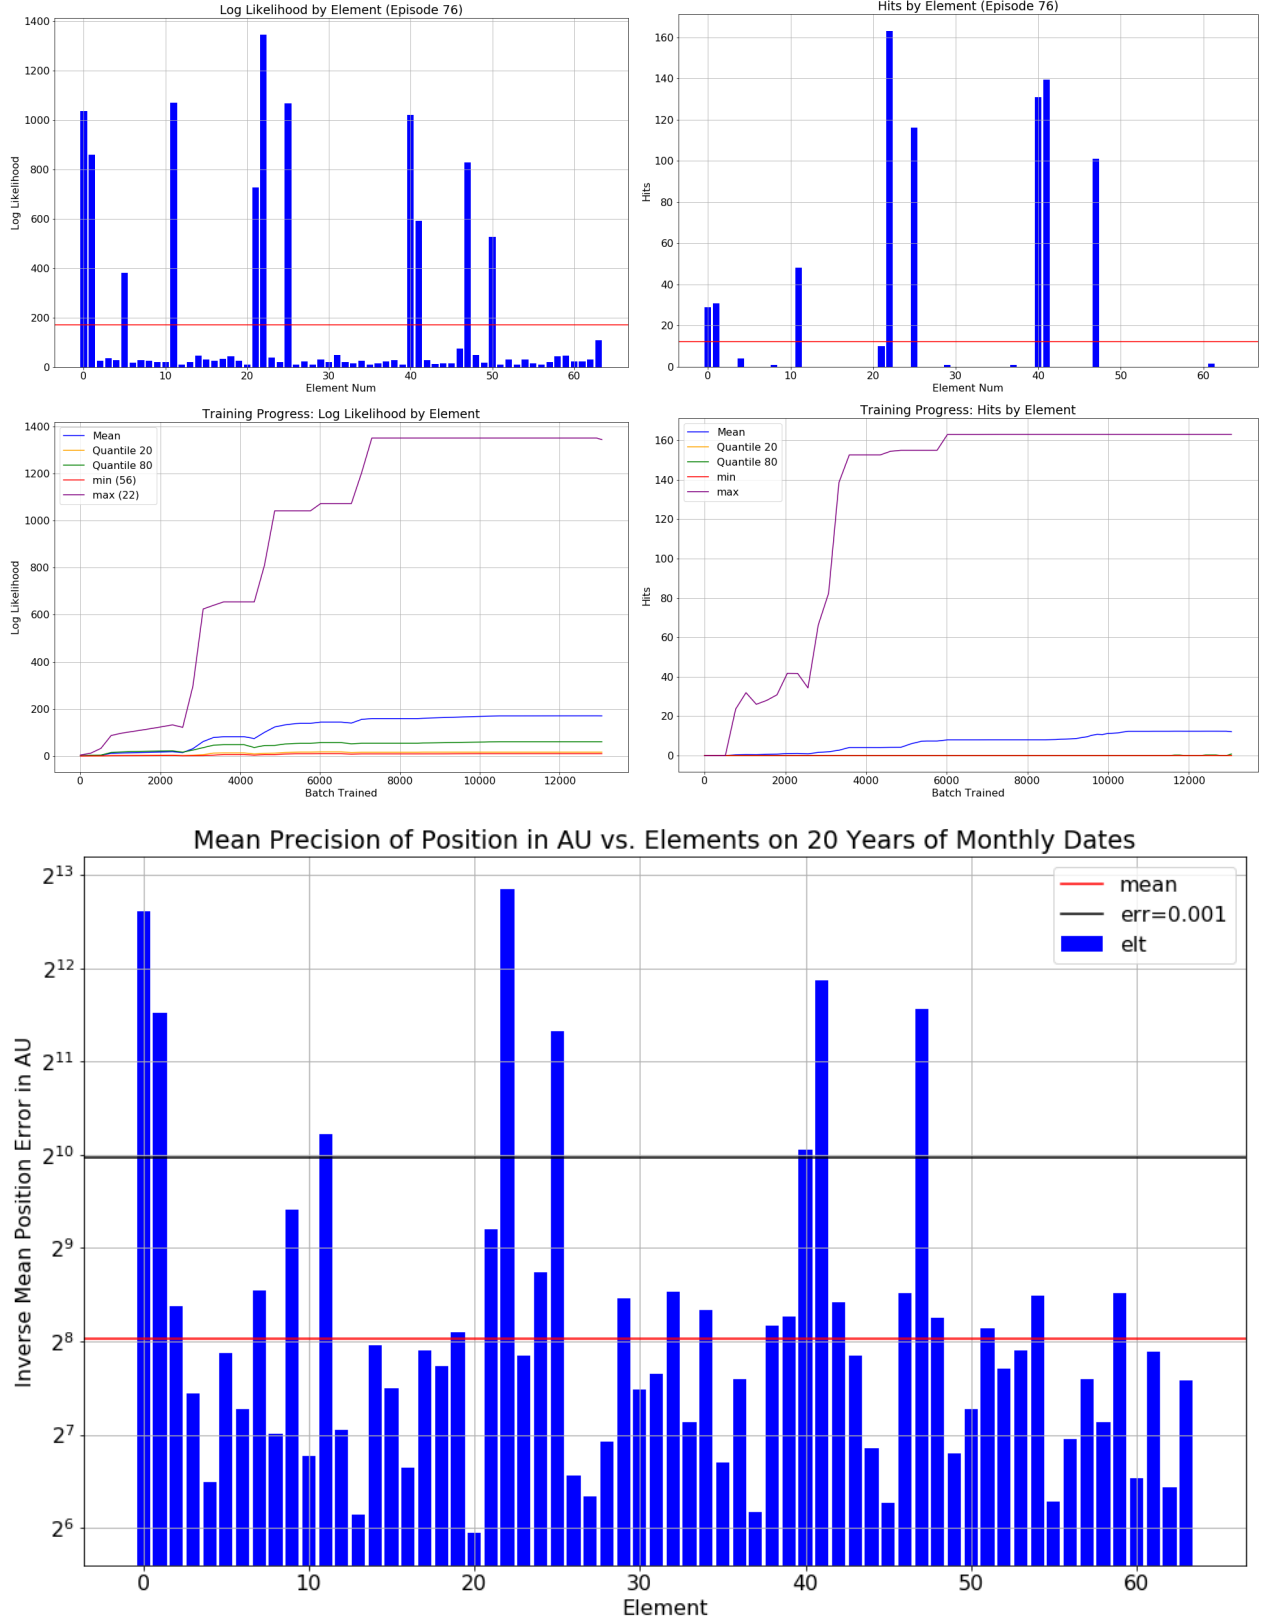


Figure 1.8: Training progress on 64 orbital elements initialized with large perturbations from real asteroids. 9 of the 64 elements converge, averaging 87 hits.

1.4 Searching for Asteroids with Random Initializations

1.5 Presenting [N] Previously Unknown Asteroids

1.6 Conclusion

1.7 Future Work

References

- [1] Hanno Rein, Shang-Fei Liu
REBOUND: An open-source multi-purpose N-body code for collisional dynamics.
Astronomy & Astrophysics. November 11, 2011.
arXiv: 1110.4876v2
- [2] Hanno Rein, David S. Spiegel
IAS15: A fast, adaptive high-order integrator for gravitational dynamics, accurate to machine precision over a billion orbits.
Monthly Notices of the Royal Astronomical Society. Printed 16 October 2014.
arXiv: 1405.4779.v2
- [3] Murray, C. D.; Dermott, S.F.
Solar System Dynamics
Cambridge University Press. 1999
- [4] Joseph Blitzstein, Jessica Hwang
Introduction to Probability
CRC Press. 2019 (Second Edition)